

FindResearch.org: How to Encourage Sharing of Research Artifacts

Christian Collberg

Todd Proebsting

Department of Computer Science
University of Arizona

<http://repeatability.cs.arizona.edu>

<http://findresearch.org>

Communications of the ACM, March 2016

Supported by *the private foundation that must not be named*

Opening Gambit

The Deception Study

Sharing Proposals

How to Share?

Pushback

Some Computer Security Paper



Really Good Computer Science Department
Really Good School

Abstract

We present a new general technique for protecting clients in distributed systems against *Remote Man-at-the-end* (R-MATE) attacks. Such attacks occur in settings where an adversary has physical access to an untrusted client device and can obtain an advantage from tampering with the hardware itself or the software it contains.

In our system, the trusted server overwhelms the untrusted client's analytical abilities by continuously and automatically generating and pushing to him diverse client code variants. The diversity subsystem employs a set of primitive code transformations that provide an ever-changing attack target for the adversary, making tampering difficult without this being detected by the server.

1. Introduction

Man-at-the-end (MATE) attacks occur in settings where an adversary has physical access to a device and compromises it by tampering with its hardware or software. *Remote man-at-the-end* (R-MATE) attacks occur in distributed systems where *untrusted clients* are in frequent communication with *trusted servers* over a network, and malicious user can get an advantage by compromising an untrusted device.

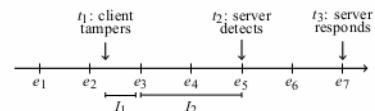
To illustrate the ubiquity of R-MATE vulnerabilities, consider the following four scenarios. First, in the *Advanced Metering Infrastructure* (AMI) for controlling the electrical power grid, networked devices ("smart meters") are installed at individual house-holds to allow two-way communication with control servers of the utility company. In an R-MATE attack against the AMI, a malicious consumer tampers with the meter to emulate an imminent blackout, or to trick a control server to send disconnect commands to other customers [7, 21]. Second, massive multiplayer online games are susceptible to R-MATE attacks since a malicious player who tampers with the game client can get an advantage over other players [16]. Third, wireless sensors are often deployed in unsecured environments (such as theaters of war) where they are vulnerable to tampering attempts. A compromised sensor could be coached into supplying the wrong observations to a base station, causing real-world damage. Finally, while electronic health records (EHR) are typically protected by encryption while stored in databases and in transit to doctors' offices, they are vulnerable to R-MATE attack if an individual doctor's client machine is compromised.

1.1 Overview

In each of the scenarios above the adversary's goal is to tamper with the client code and data under his control. The trusted server's goal is to *detect* any such integrity violations, after which countermeasures (such as severing connections, legal remedies, etc.) can be launched.

Security mechanisms. In this paper we present a system that achieves protection against R-MATE attacks through the extensive use of code diversity and continuous code replacement. In our system, the trusted server continuously and automatically generates diverse variants of client code, pushes these code updates to the untrusted clients, and installs them as the client is running. The intention is to force the client to constantly analyze and re-analyze incoming code variants, thereby overwhelming his analytical abilities, and making it difficult for him to tamper with the continuously changing code without this being detected by the trusted server.

Limitations. Our system specifically targets distributed applications which have frequent client-server communication, since client tampering can only be detected at client-server interaction events. Furthermore, while our use of code diversity can *delay* an attack, it cannot completely *prevent* it. Our goal is therefore the rapid *detection* of attacks; applications which need to completely prevent any tampering of client code, for even the shortest length of time, are not suitable targets for our system. To see this, consider the following timeline in the history of a distributed application running under our system:



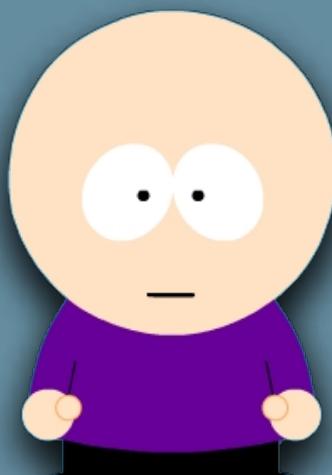
The e_i 's are *interaction events*, points in time when clients communicate with servers either to exchange application data or to perform code updates. At time t_1 the client tampers with the code under his control. Until the next interaction event, during interval I_1 , the client runs autonomously, and the server cannot detect the attack. At time t_2 , after an interval I_2 consisting of a few interaction events, the client's tampering has caused it to display anomalous behavior, perhaps through the use of an outdated communication protocol, and the server detects this. At time t_3 , finally, the server issues a response, perhaps by shutting

Really
famous
authors



To: authors@cs.ux.edu

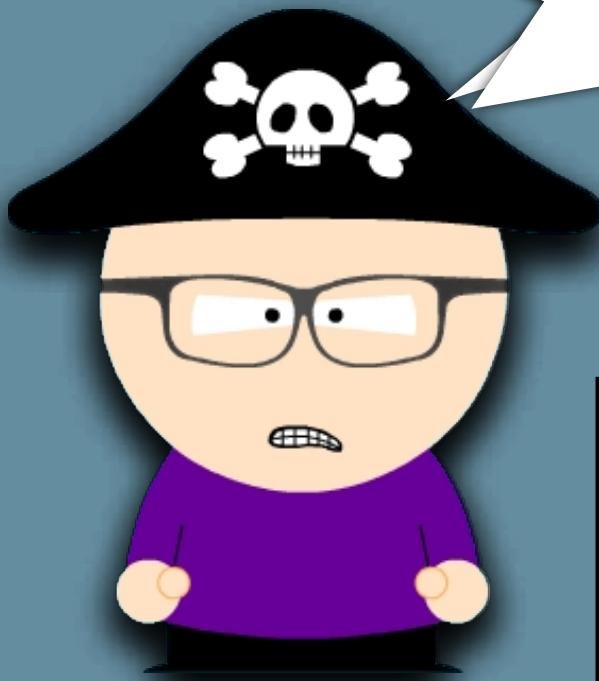
Cool paper! Can you send
me the system so I can
break it? 😊



Technical
Report

Conference
Paper

PhD
Thesis



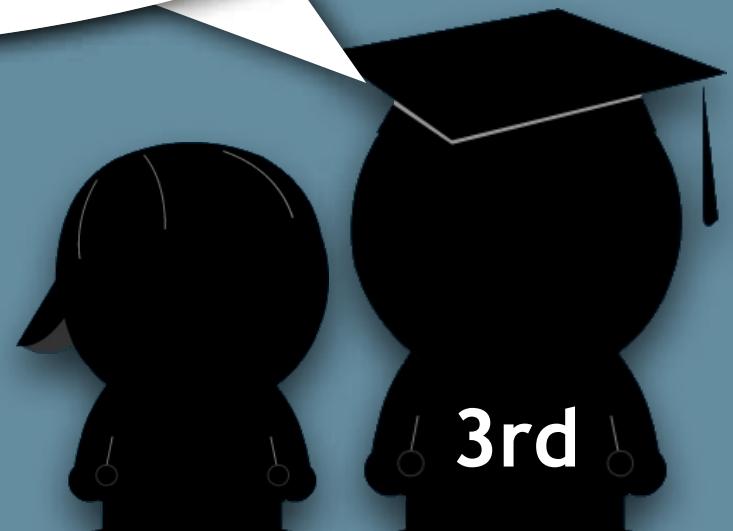
My version of
their software!

f: never used!
g: not defined!
h: makes no sense!

```
type operator =  
| A  
| B of operand * value * binop  
| C of operand * value * operand * binop  
| D of operand * value * operand * binop  
| E of operand * operand
```

To: authors@cs.ux.edu

- 1) Why is f unused?
- 2) Define g, please!
- 3) Explain h, please!



To: PI,DC@cs.ux.edu

I ... request under the
OPEN RECORDS ACT ... ALL
RESEARCH ARTIFACTS ...

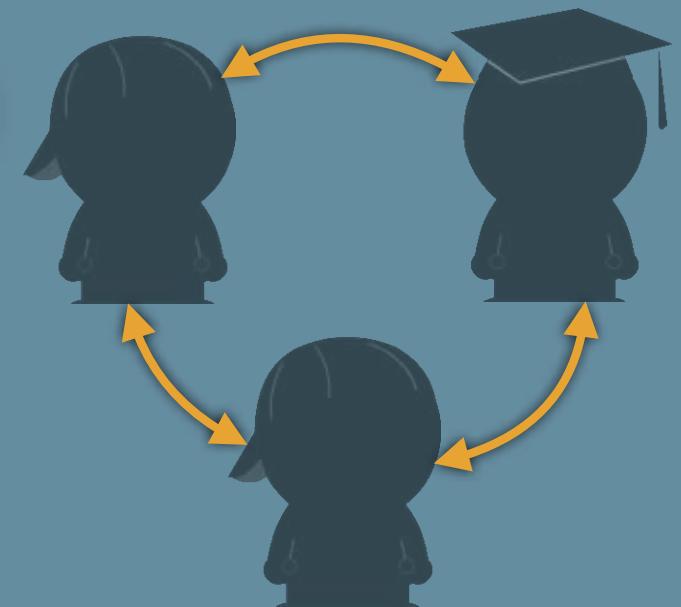
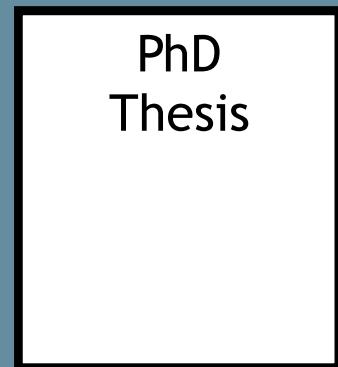
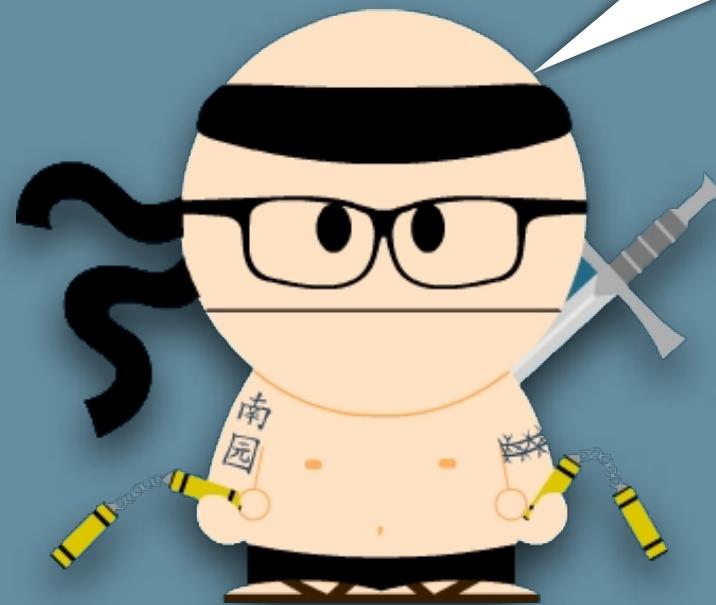


From: legal@cs.ux.edu

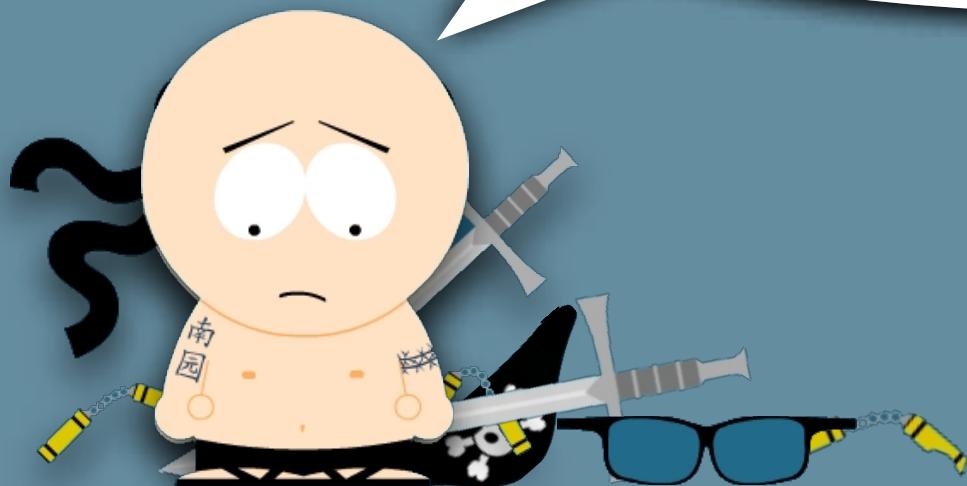
... to the extent such records may exist, **they will not be produced pursuant to ORA.**



You
can't find it?
Seriously?



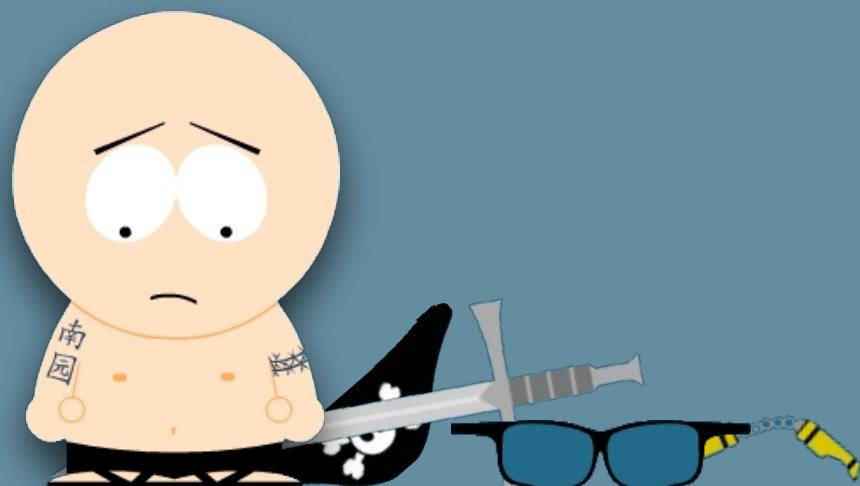
... we estimate a total cost of **\$2,263.66** to search for, retrieve, redact and produce such records.





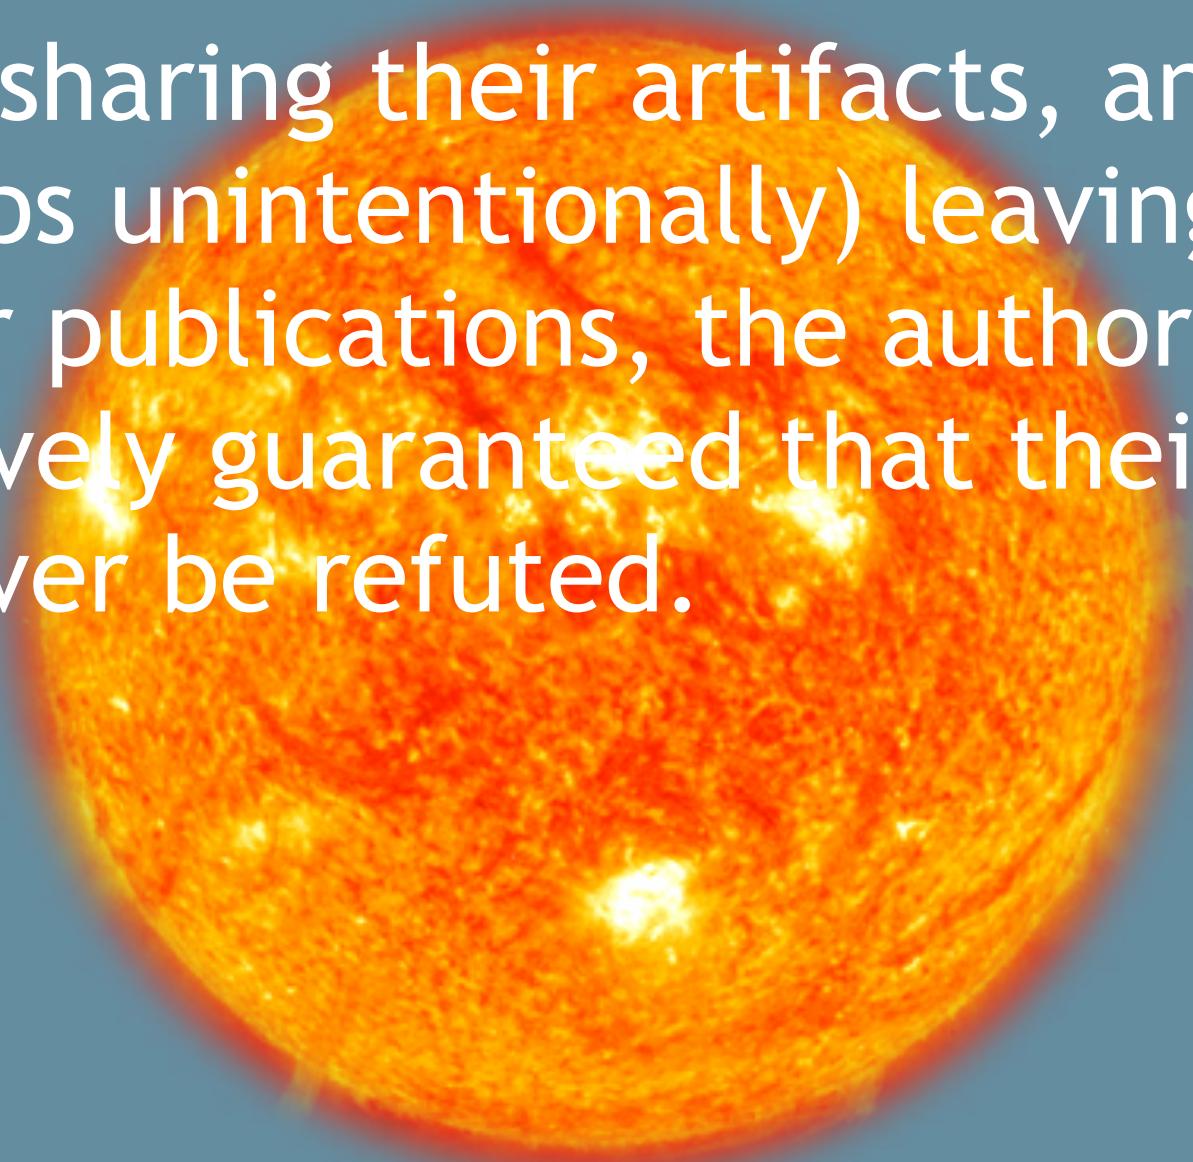
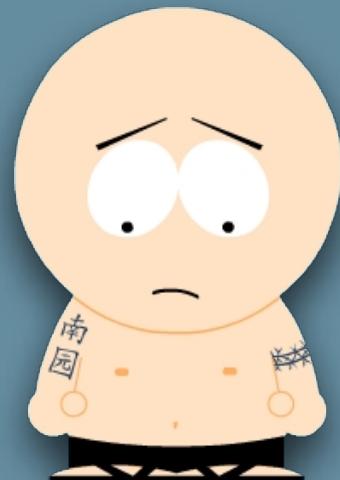
Grant application
#: [REDACTED]

We will also make our data
and software available to
the research community
when appropriate.



Consequences

By not sharing their artifacts, and by (perhaps unintentionally) leaving holes in their publications, the authors have effectively guaranteed that their claims can never be refuted.

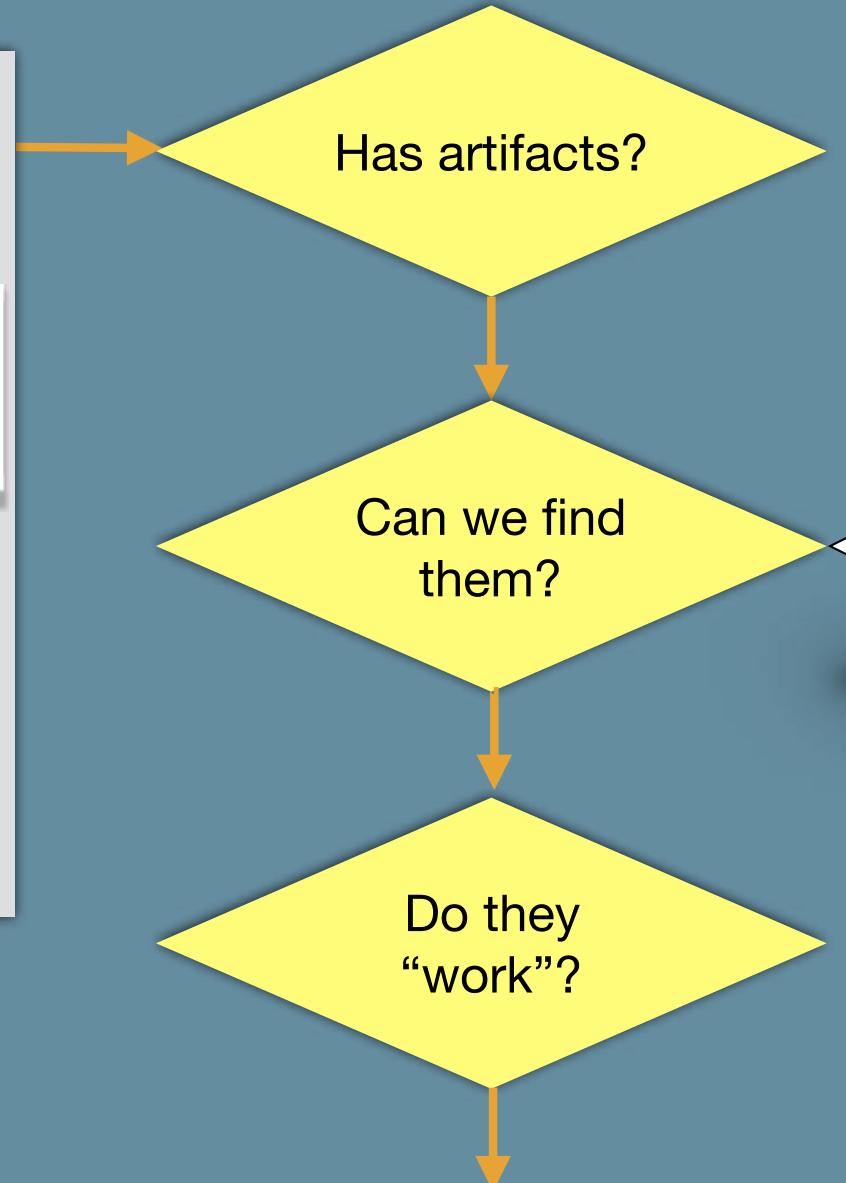


1st Law of Artifact Sharing

The probability of getting code out of someone is inversely proportional to the outrageousness of the claims in the paper.

The Deception Study

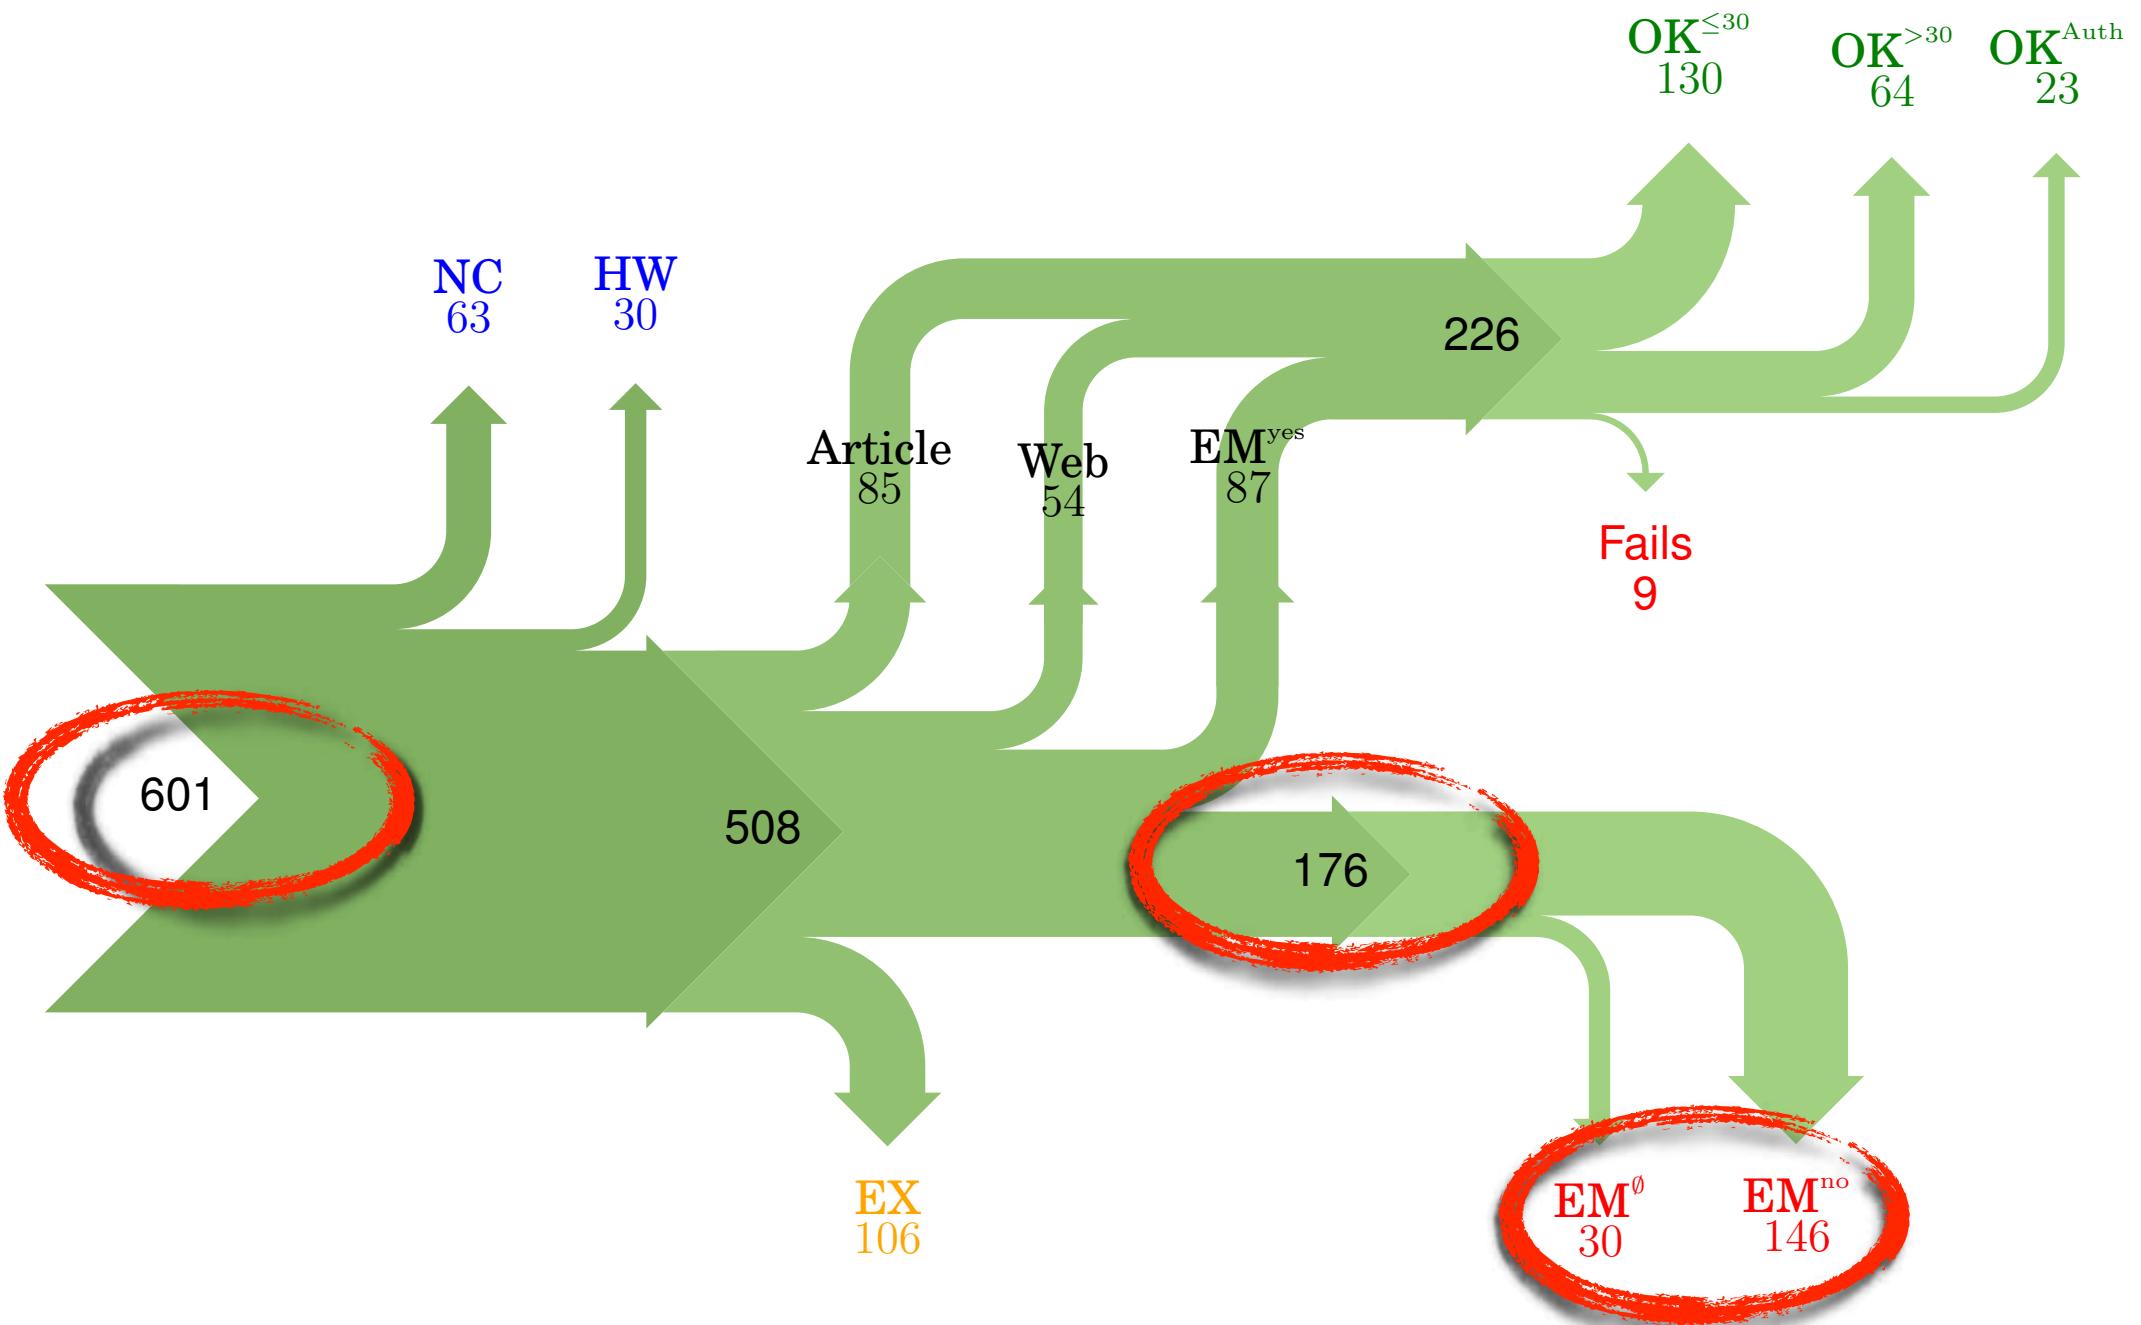
601 Research Papers



1. Article?
2. Web?
3. Email?

Weakly Repeatable

Authors make the artifacts used to create the results in their article available, and they build.





The good news ... I was able to find some code. I am just **hoping** that it ... **matches the implementation** we ... used for the paper.

Versioning



The code was **never intended to be released** so is not in any shape for general use.

No Intention to Share

[Our] prototype ... included many moving pieces that only [student] knew how to operate ... **he left.**



Personnel Issues

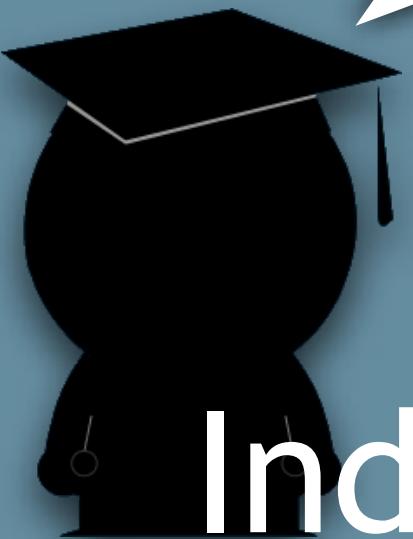
... the server in which my
implementation was
stored had a **disk crash**
... three disks crashed ...
Sorry for that.



Lost Code

[Therefore] we will not provide the software outside the group.



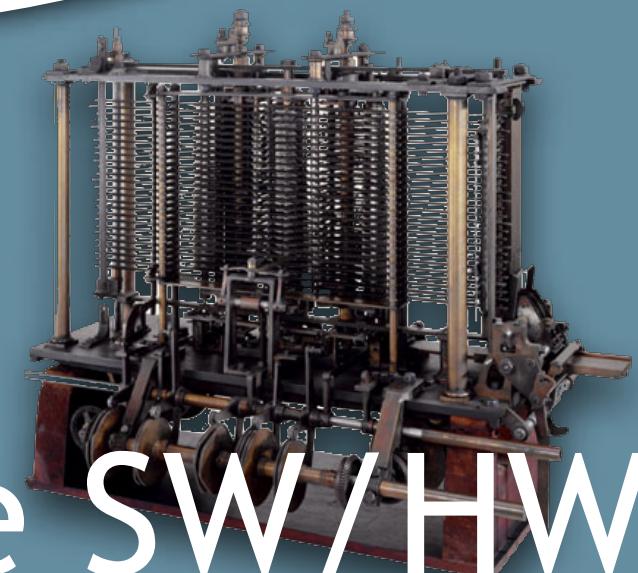


**... we can't share what
did for this paper. ...**
this is not in the
academic tradition, but
this is a hazard in an
industrial lab.

Industrial Lab Tradeoffs

... we have no plans to make the scheduler's source code publicly available. ... because [ancient OS] as such

does not exist anymore



Obsolete SW/HW



... we have an agreement
with the [business], and
we cannot release the
code because of the
potential **privacy risks**

...

Privacy/Security

Sharing Proposals

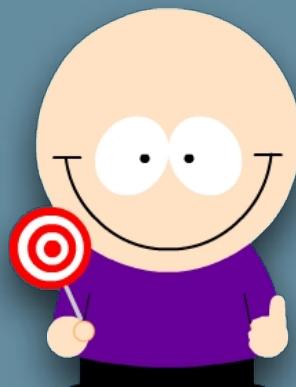
1. Artifact Evaluation
2. Artifact Repositories
3. Funding Agency Audits
4. Sharing Specifications

Artifact Evaluation



Artifact evaluation is open only to accepted papers. This is intentional: it ensures that the Artifact Evaluation Committee **cannot influence whether or not a paper is accepted**.

- 60 accepted papers
- 27 (45%) met or exceeded expectations



Dates

Paper decision notification:
Feb 5, 2014
Artifacts due:
Feb 10, 2014
Decisions announced:
approx. Mar 15, 2014
Camera-ready due:
Mar 20, 2014

Packaging

How to Submit

Please read the [guidelines](#) on what to submit.
Please upload your submission to [EasyChair](#).

The Committee

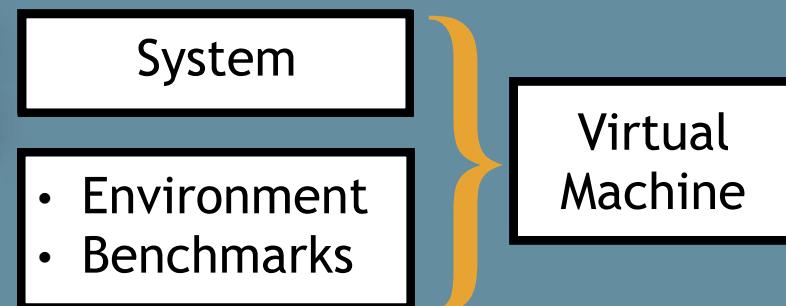
The committee consists of several up-and-coming researchers with [Eric Eide](#), [Shriram Krishnamurthi](#), and [Jan Vitek](#) heading the process.

Process

Artifact evaluation is open only to accepted papers. This is intentional: it ensures that the AEC cannot influence whether or not a paper is accepted. This measure was put in place to reassure authors who felt this would be too radical a change to the process of evaluating conference paper submissions.

Of course, this doesn't mean you can't start getting ready! We have published the

Artifact Repositories



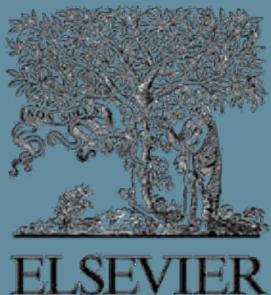
- If you build it, they still won't come...

Funding Agency Audits



Agencies should conduct **random audits** to ensure that research artifacts are shared in accordance with what was promised in the grant application

Sharing Specifications



Title	
Abstract	Introduction
.....
.....
.....
.....
Keywords
.....
Copyright	
.....	Sharing
.....
.....

Provided both in

- submission
- final paper

Clarify which research artifacts will be available

Publishers should require articles to contain a **sharing contract** specifying the level of repeatability to which its authors will commit

Sharing Proposals

5. Artifact Indexing

FindResearch.org

Longitudinal Study

Congrats on
your new paper!

Will you share?

Sure! Nope!



Public Sharing web site



ACM Programming Language Design and Implementation, PLDI 2014		PLDI 2014
Proposed	Received	Accepted
Algorithmic primitives for distributed systems		Algorithmic primitives for distributed systems
Java 8		Java 8
Modern, memory-managing, concurrent programs in C/C++		Modern, memory-managing, concurrent programs in C/C++
Object-Oriented Software Engineering: From Object-Oriented Analysis and Design to Object-Oriented Systems and Tools		Object-Oriented Software Engineering: From Object-Oriented Analysis and Design to Object-Oriented Systems and Tools
Programmatic access points in communication	• Implementing access points in WebAssembly	Programmatic access points in communication
Smart Cities, Smart Cities	• Implementing access points in WebAssembly	Smart Cities, Smart Cities
Modernizing Software Engineering		Modernizing Software Engineering
Mobile, Mobile, Mobile		Mobile, Mobile, Mobile
Designing the future of wireless networks and mobile computing	• Implementing access points in WebAssembly	Designing the future of wireless networks and mobile computing
Level of programming		Level of programming
Tool support		Tool support
Programmatic access to object-oriented systems	• Implementing access points in WebAssembly	Programmatic access to object-oriented systems
Tool integration, language integration, domain integration, API integration	• Implementing access points in WebAssembly	Tool integration, language integration, domain integration, API integration
Tool integration, domain integration, API integration		Tool integration, domain integration, API integration

100 conferences

1. Mo
1000s of authors,
2. D

10,000s of papers, 2. Dir

over 5 years!³

- 1. Motivating researchers to share their papers, sharing, co-authors, license, dis-
 - 2. Directory of research artifacts URL,
 - 3. Trending data for funding agencies covariate,
 - 4. Data for repeatability research... funding...

Sharing, community rs to share license, discussion!

↳ [LICENCIATE](#) discussion!

A Catalog of Research Artifacts for Computer Science

FindResearch.org aims
(e.g., code and data)
essential for repeating

Se

E.g. Au

- 79 conferences
- 8721 articles
- 20,675 unique authors
- 28,097 survey emails sent

Complete results pending while our site is being indexed

Find research artifacts by publication venue

Publication Venue

Year (#artifacts)

ACM Architectural Support for Programming Languages and Operating Systems, ASPLOS	2016 (20)
ACM Computer and Communications Security, CCS	2016 (68) 2015 (52) 2014 (22)
ACM Conference on Mobile Systems, Applications, and Services, MobiSys	2016 (8)
ACM Hypertext and Social Media, HT	2016 (6) 2015 (3) 2014 (9)
ACM Information Retrieval, SIGIR	2016 (62) 2015 (17)
ACM International Conference on Embedded Software, EMSOFT 2016	2016 (2)
ACM International Conference on Management of Data, SIGMOD	2016 (41) 2015 (27)
ACM Internet Measurement Conference, IMC	2016 (16)

ACM Principles of Programming Languages, POPL 2017

Title/Authors	Research Artifacts [?]	Details
<i>Serializability for eventual consistency: criterion, analysis, and applications</i> Lucas Brutschy, Dimitar Dimitrov, Peter Müller, Martin T. Vechev	<ul style="list-style-type: none">http://ecracer.inf.ethz.ch/ 	Author Comments: Discussion Comments: 0 Sharing: Research produced artifacts Verification: Authors have verified information More...
<i>Deciding equivalence with sums and the empty type</i> Gabriel S.		Discussion Comments: 0 Verification: Author has not verified information More...
<i>A semantics for probabilistic programs</i> Arthur Azriel, Shin-ya Kuroda		Author Comments: Discussion Comments: 0 Sharing: Research produced no artifacts Verification: Authors have verified information More...
<i>Exact Bayesian inference for probabilistic programs</i> Chung-chuan Chen, Norman Ramsey		Author Comments: Discussion Comments: 0 Sharing: Research produced artifacts Verification: Authors have verified information More...
<i>Intersection type calculi of bounded dimension</i> Andreas Dudenhefner, Jakob Rehof		Discussion Comments: 0 Verification: Authors have not verified information More...

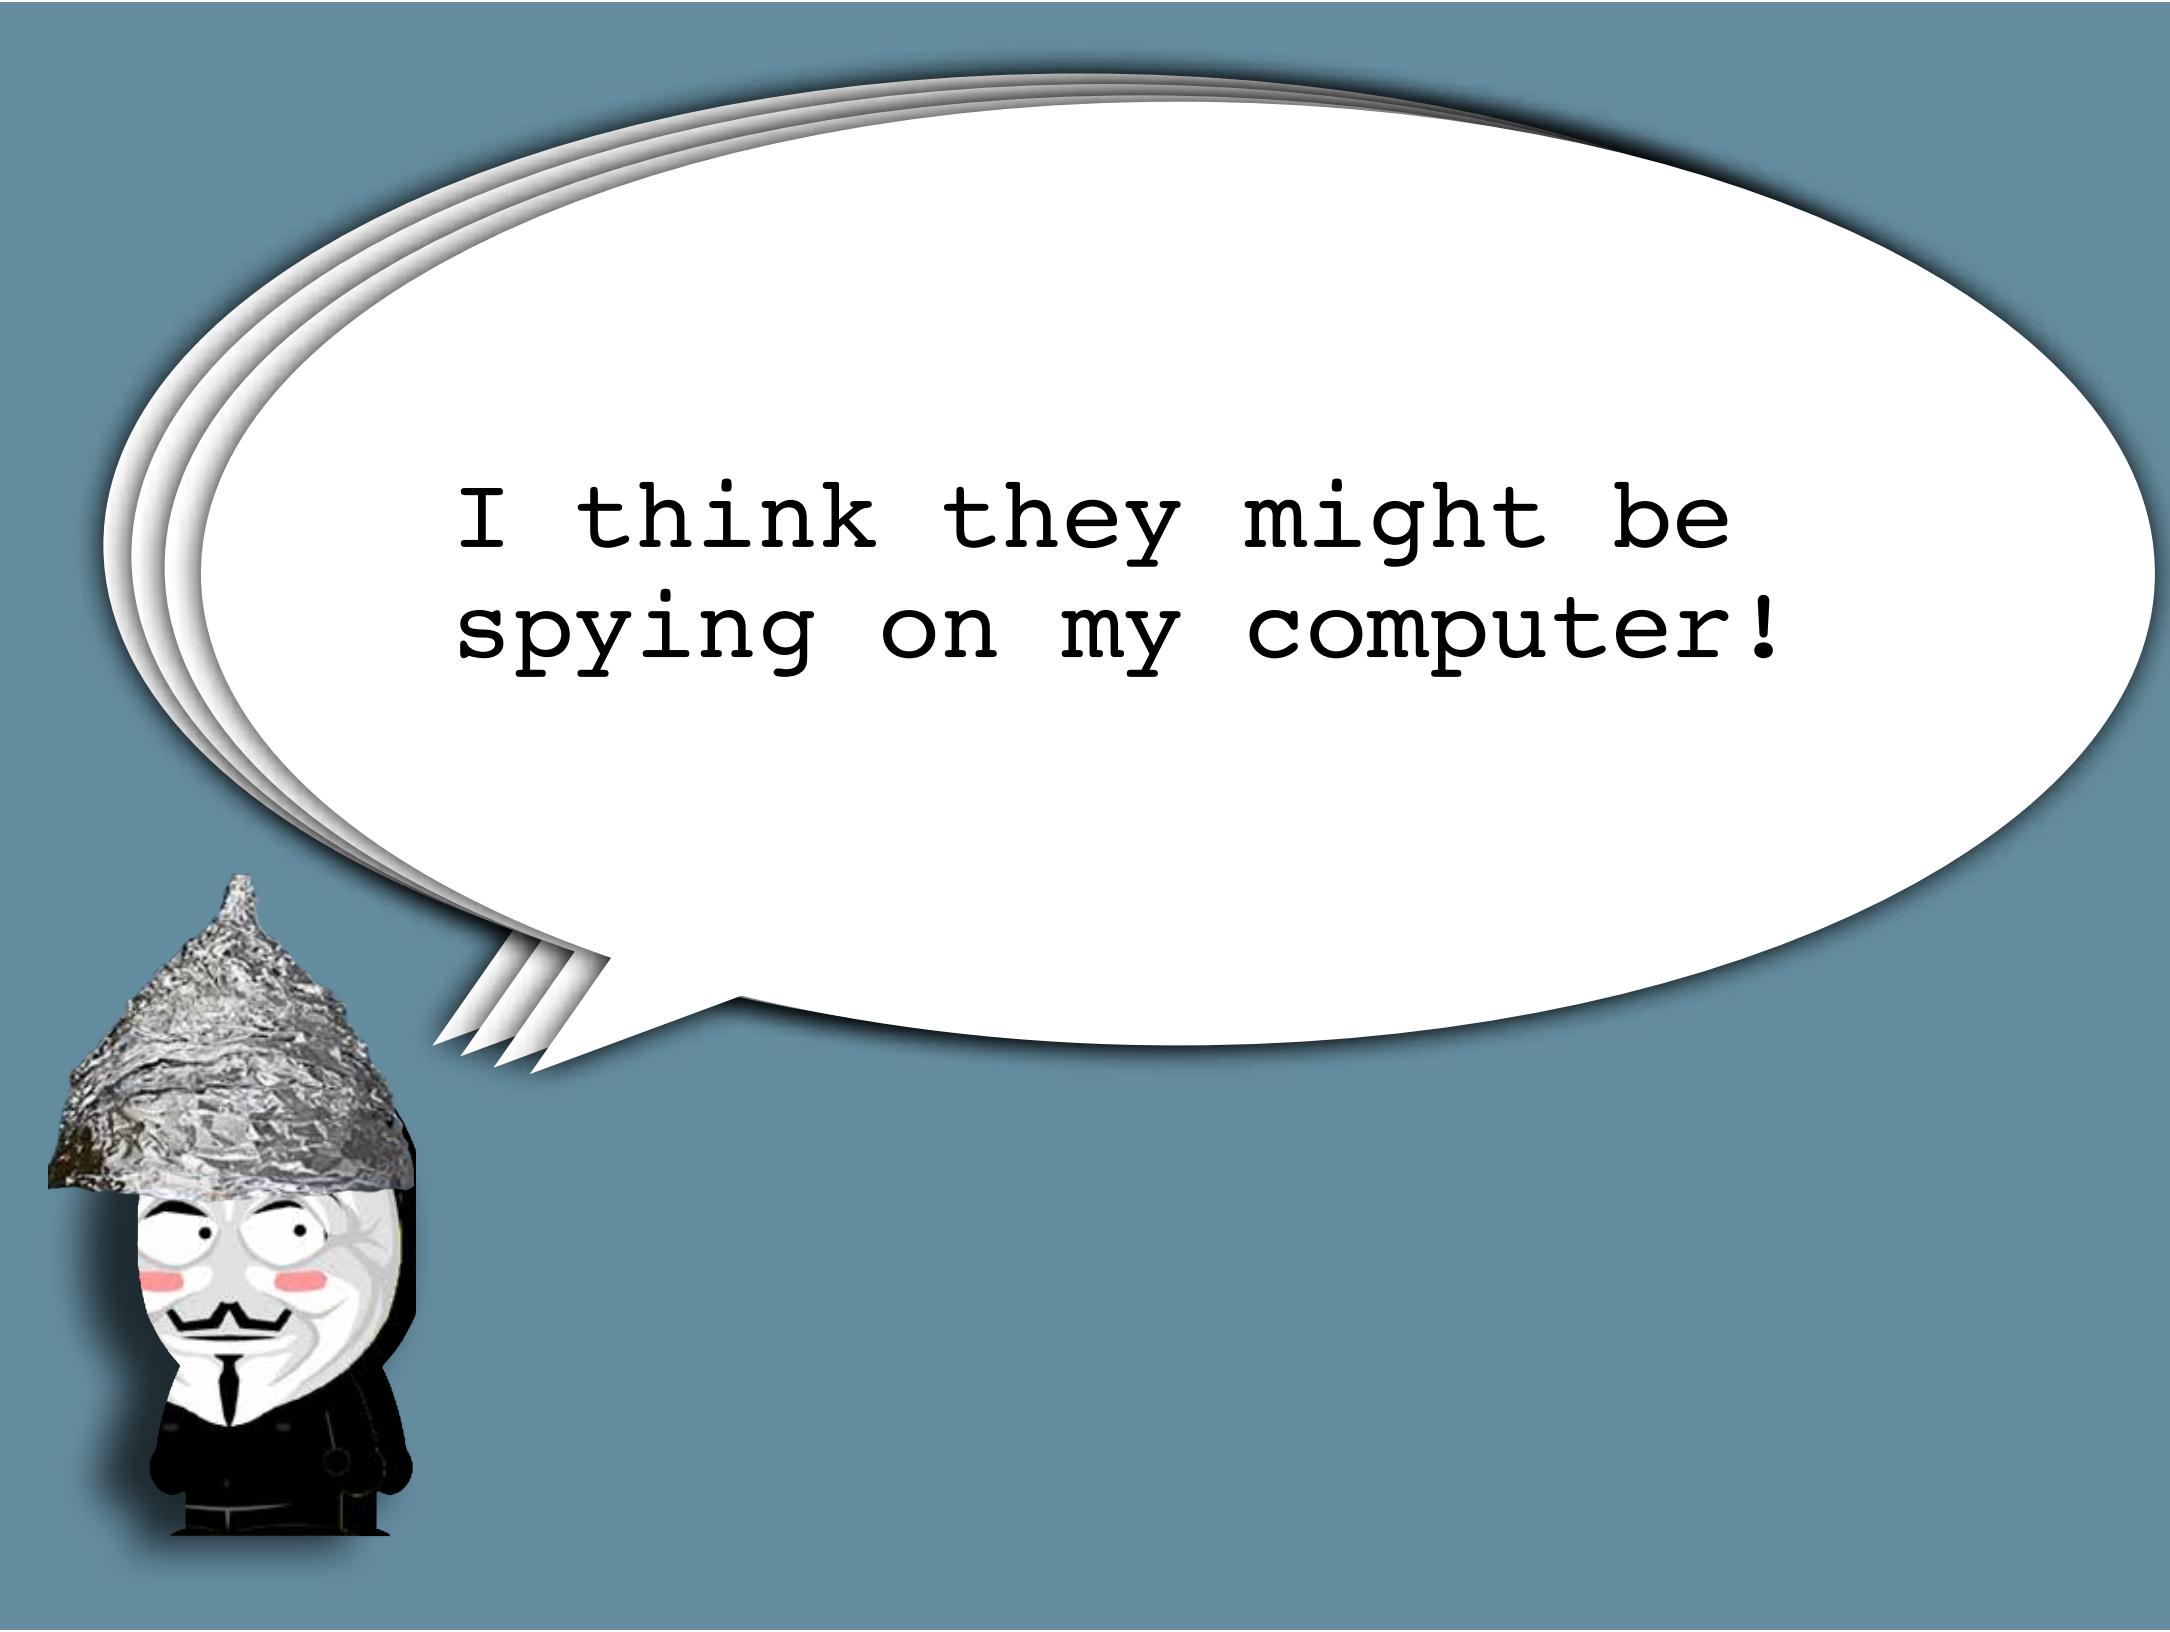
- Verified articles: 813 (9%)
- Articles with shared artifacts: 521 (6%)
- Unable to share: 108



This is a great project -
keep up the good work,
folks!



We'd like to help by
providing the information
you need for our
conference!



I think they might be
spying on my computer!



How to Share?

- What to Share?
- How to Package?
- Where to Host?

Your colleagues may need access to an artifact for many different reasons:

- They may want to read the source code to better understand your paper.
- They may want to rerun the experiments you present in your paper to ensure you got things right.
- They may want to build upon and extend your work.
- They may want to compare your results to their own.
- They may want to run different experiments on your code, or run the same experiments you ran but on different data sets.

To ensure *repeatability* provide all the

- sources
- makefiles
- external libraries
- instructions for proof assistants
- data sets
- installation instructions
- scripts to run experiments

that went into producing the results reported on in the final, published, paper.

To ensure *longevity*, future-proof by including all libraries.

Install the following dependencies:

```
sudo apt-get -y install git g++ cmake libboost-regex-dev
```

Will **apt-get abclib/1.2.3** work in 5 years?

To ensure identical binaries are built,
precisely document software you can't
include:

- Name: abclib
- Version: 1.2.3
- Location: <http://abclib.org>

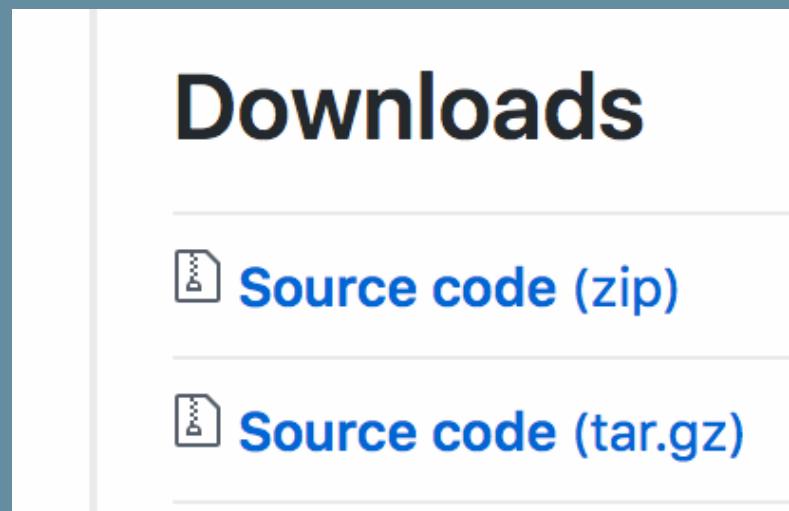
Make sure you have installed on your machine

- git
- a recent version of GCC (version 4.7 or higher)
- flex (version 2.5 or higher)

Beware of overloaded package names!

To ensure a *clear mapping* between paper and artifact, provide

- A permanent package (zip-file, virtual machine, container), or
- A "tagged" version of a public repository



2nd Law of Artifact Sharing

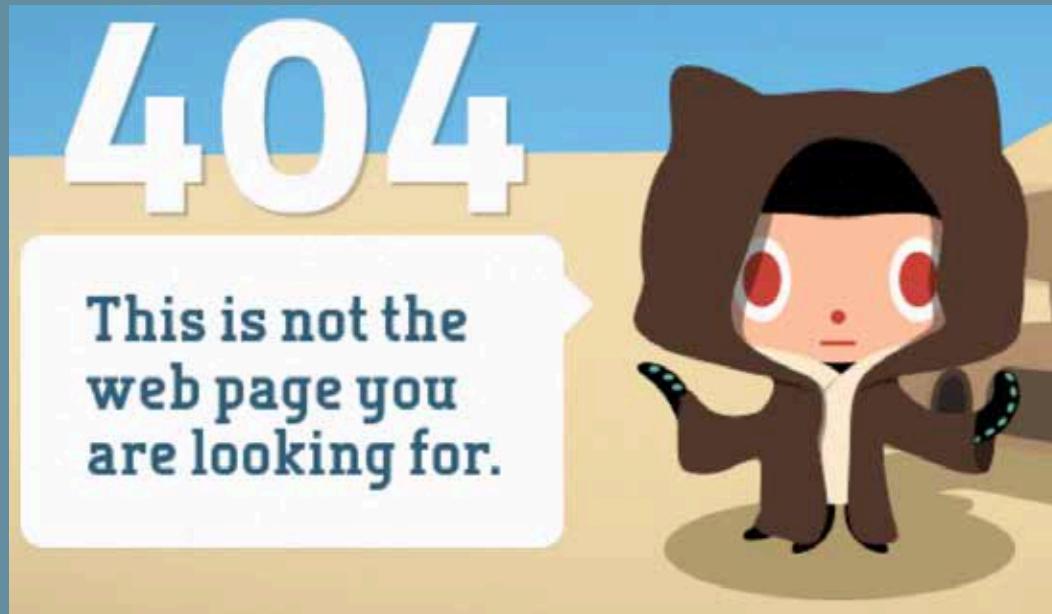
If you can't find which artifact goes with the final published version of a paper, if you don't have every single bit of that artifact, if you don't know the exact environment in which experiments were run, you ain't got nothing.

To ensure availability, host on github, amazon, azure, ...

Service Temporarily Unavailable

The server is temporarily unable to service your request due to maintenance downtime or capacity problems. Please try again later.

Apache Server at eventbuilder.geovid.org Port 80



18 verified papers with shared artifacts appear to have broken links (3%)

Some Computer Security Paper



Really Good Computer Science Department
Really Good School

Abstract

We present a new general technique for protecting clients in distributed systems against *Remote Man-in-the-end* (R-MATE) attacks. Such attacks occur in settings where an adversary has physical access to an untrusted client device and can obtain an advantage from tampering with the hardware itself or the software it contains.

In our system, the trusted server overwhelms the untrusted client's analytical abilities by continuously and automatically generating and pushing to him diverse client code variants. The diversity subsystem employs a set of primitive code transformations that provide an ever-changing attack target for the adversary, making tampering difficult without this being detected by the server.

1. Introduction

Man-in-the-end (MATE) attacks occur in settings where an adversary has physical access to a device and compromises it by tampering with its hardware or software. *Remote man-in-the-end* (R-MATE) attacks occur in distributed systems where *untrusted clients* are in frequent communication with *trusted servers* over a network, and malicious user can get an advantage by compromising an untrusted device.

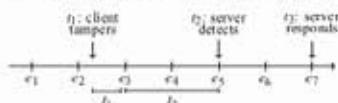
To illustrate the ubiquity of R-MATE vulnerabilities, consider the following four scenarios. First, in the *Advanced Metering Infrastructure* (AMI) for controlling the electrical power grid, networked devices ("smart meters") are installed at individual households to allow two-way communication with control servers of the utility company. In an R-MATE attack against the AMI, a malicious consumer tampers with the meter to emulate an imminent blackout, or to trick a control server to send disconnect commands to other customers [2][2]. Second, massive multiplayer online games are susceptible to R-MATE attacks since a malicious player who tampers with the game client can get an advantage over other players [16]. Third, wireless sensors are often deployed in unsecured environments (such as theaters of war) where they are vulnerable to tampering attempts. A compromised sensor could be coaxed into supplying the wrong observations to a base station, causing real-world damage. Finally, while electronic health records (EHR) are typically protected by encryption while stored in databases and in transit to doctors' offices, they are vulnerable to R-MATE attack if an individual doctor's client machine is compromised.

1.1 Overview

In each of the scenarios above the adversary's goal is to tamper with the client code and data under his control. The trusted server's goal is to *detect* any such integrity violations, after which countermeasures (such as severing connections, legal remedies, etc.) can be launched.

Security mechanisms. In this paper we present a system that achieves protection against R-MATE attacks through the extensive use of code diversity and continuous code replacement. In our system, the trusted server continuously and automatically generates diverse variants of client code, pushes these code updates to the untrusted clients, and installs them as the client is running. The intention is to force the client to constantly analyze and re-analyze incoming code variants, thereby overwhelming his analytical abilities, and making it difficult for him to tamper with the continuously changing code without this being detected by the trusted server.

Limitations. Our system specifically targets distributed applications which have frequent client-server communication, since client tampering can only be detected at client-server interaction events. Furthermore, while our use of code diversity can *delay* an attack, it cannot completely *prevent* it. Our goal is therefore the rapid *detection* of attacks; applications which need to completely prevent any tampering of client code, for even the shortest lengths of time, are not suitable targets for our system. To see this, consider the following timeline in the history of a distributed application running under our system:



The e_i 's are *interaction events*, points in time when clients communicate with servers either to exchange application data or to perform code updates. At time t_1 the client tampers with the code under his control. Until the next interaction event, during interval I_1 , the client runs autonomously, and the server cannot detect the attack. At time t_2 , after an interval I_1 consisting of a few interaction events, the client's tampering has caused it to display anomalous behavior, perhaps through the use of an outdated communication protocol, and the server detects this. At time t_3 , finally, the server issues a response, perhaps by shutting

To ensure colleagues can ask questions use permanent email addresses

- 2,591 emails bounced (9%)
- 5440 authors without an email address (21%)
- 854 articles without any email addresses (13%)

Pushback



My code is on github, what
more do you want???

Sharing ≠ Repeatability



It's impossibly hard to make a perfect package to share, so why bother?

The perfect is the enemy of the good



If you force sharing for publication, industrial labs can't publish, so let's make sharing optional.

Optional sharing



I really don't have time
for this - I need to

- publish the next paper
- ensure my students get jobs
- submit the next grant

The Dean can't read
but the Dean can count

3rd Law of Artifact Sharing

The root of the scientific reproducibility problem is sociological, not technological: we do not produce solid artifacts or attempt to replicate the work of our peers because there is little professional glory to be gained from doing so.



Some professor will convince a naive first-year student to use your code as a starting point for their experiments and you will have contributed to that poor student's misery and anguish.

Publishing code is a rather thankless task

Still, it's in the best traditions of science to give our colleagues all the ammunition they need to falsify our results. Thus, **science \leftrightarrow sharing**.



On the Importance of being Transparent

Questions?

