

# Quantitative Network Monitoring with NetQRE

Yifei Yuan<sup>1,2</sup>, Dong Lin<sup>1,3</sup>, Ankit Mishra<sup>1</sup>, Sajal Marwaha<sup>1</sup>,  
Rajeev Alur<sup>1</sup>, and Boon Thau Loo<sup>1</sup>


<sup>1</sup>University of Pennsylvania

<sup>2</sup>Carnegie Mellon University

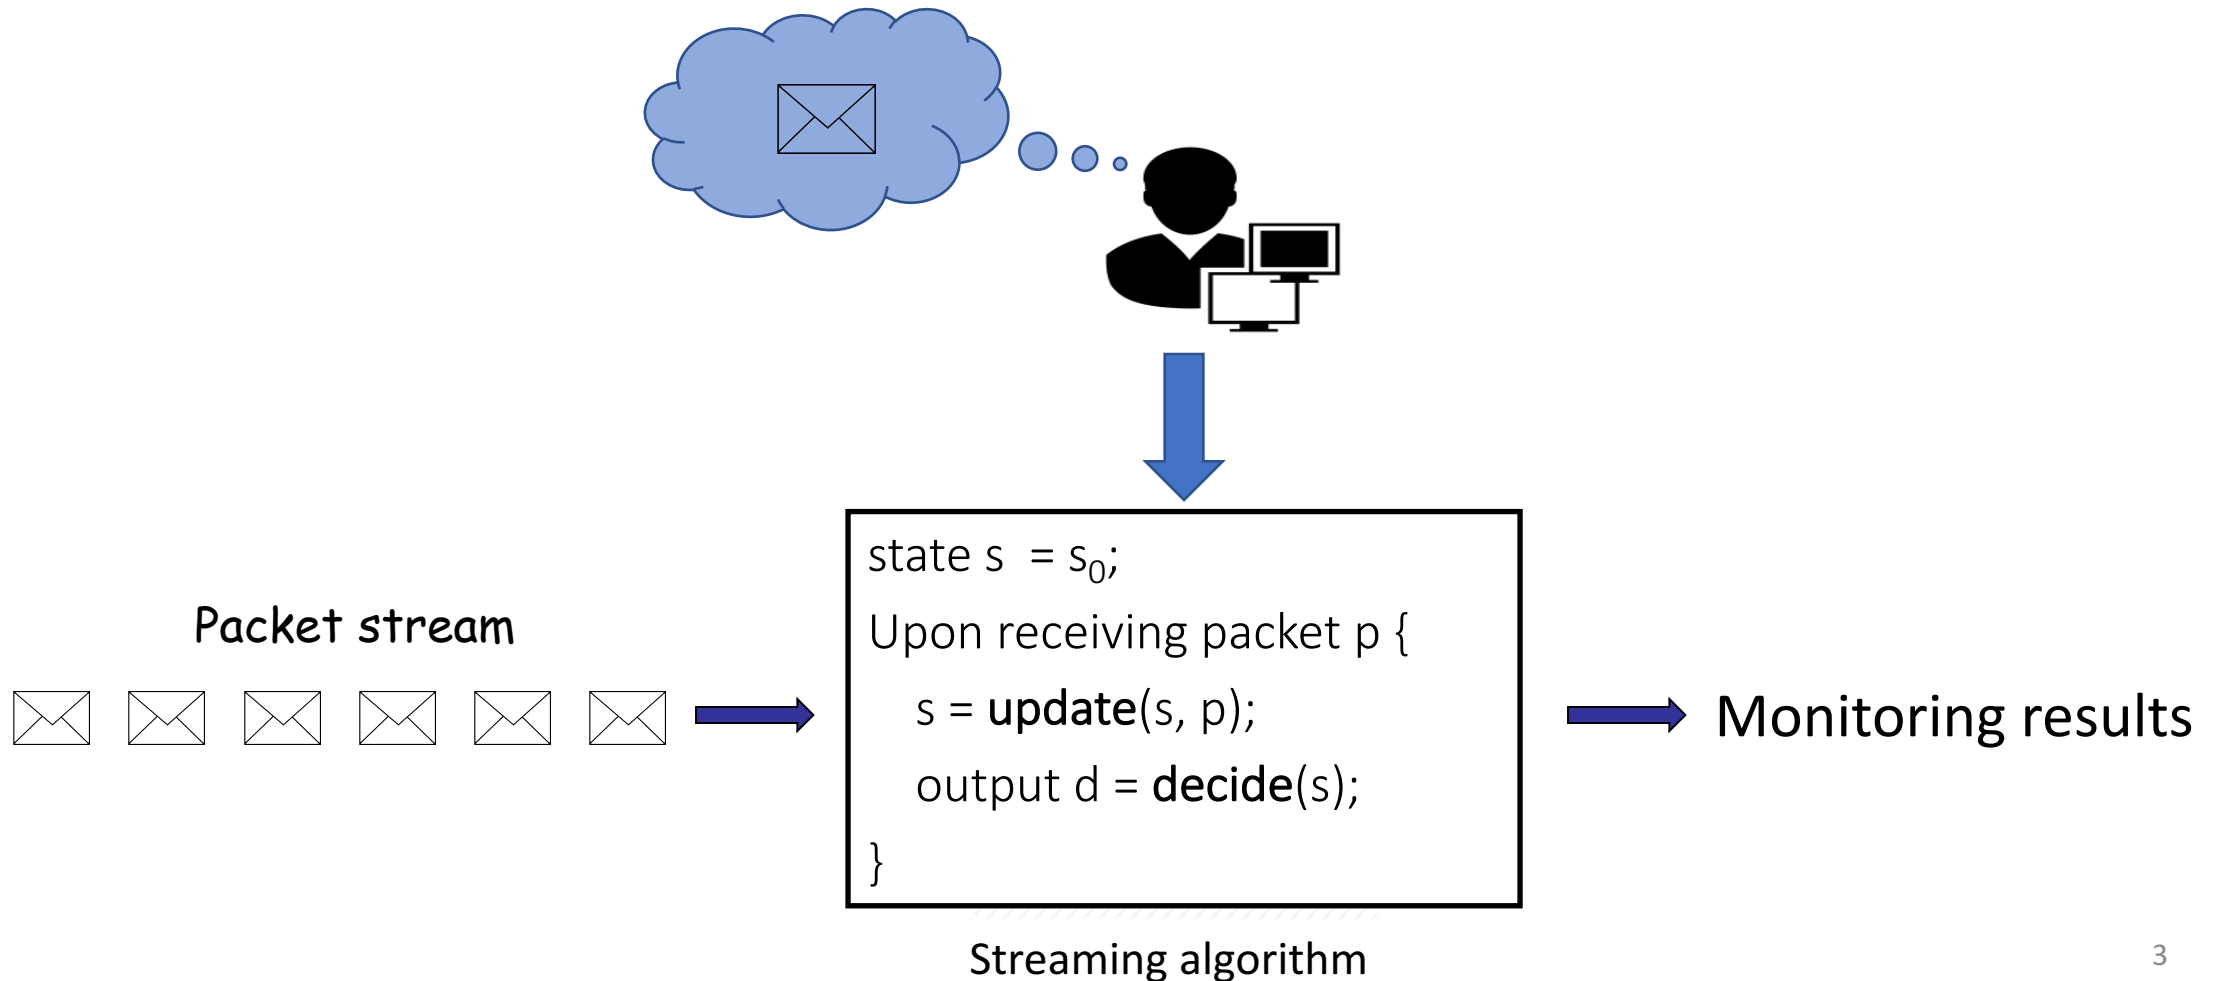
<sup>3</sup>LinkedIn Inc.

# Network Monitoring is Important


- Security
  - Heavy hitter
  - Super spreader
  - Syn flood
  - Slowloris
  - ...
- Performance
  - Traffic matrix
  - Application usage
  - ...



# Today's Low-level Programming Abstraction



# Motivating Example: VoIP Monitoring



## Example Policy:

1. Monitor average number of VoIP calls per user
2. Alert a user, if her/his number of calls exceeds a threshold

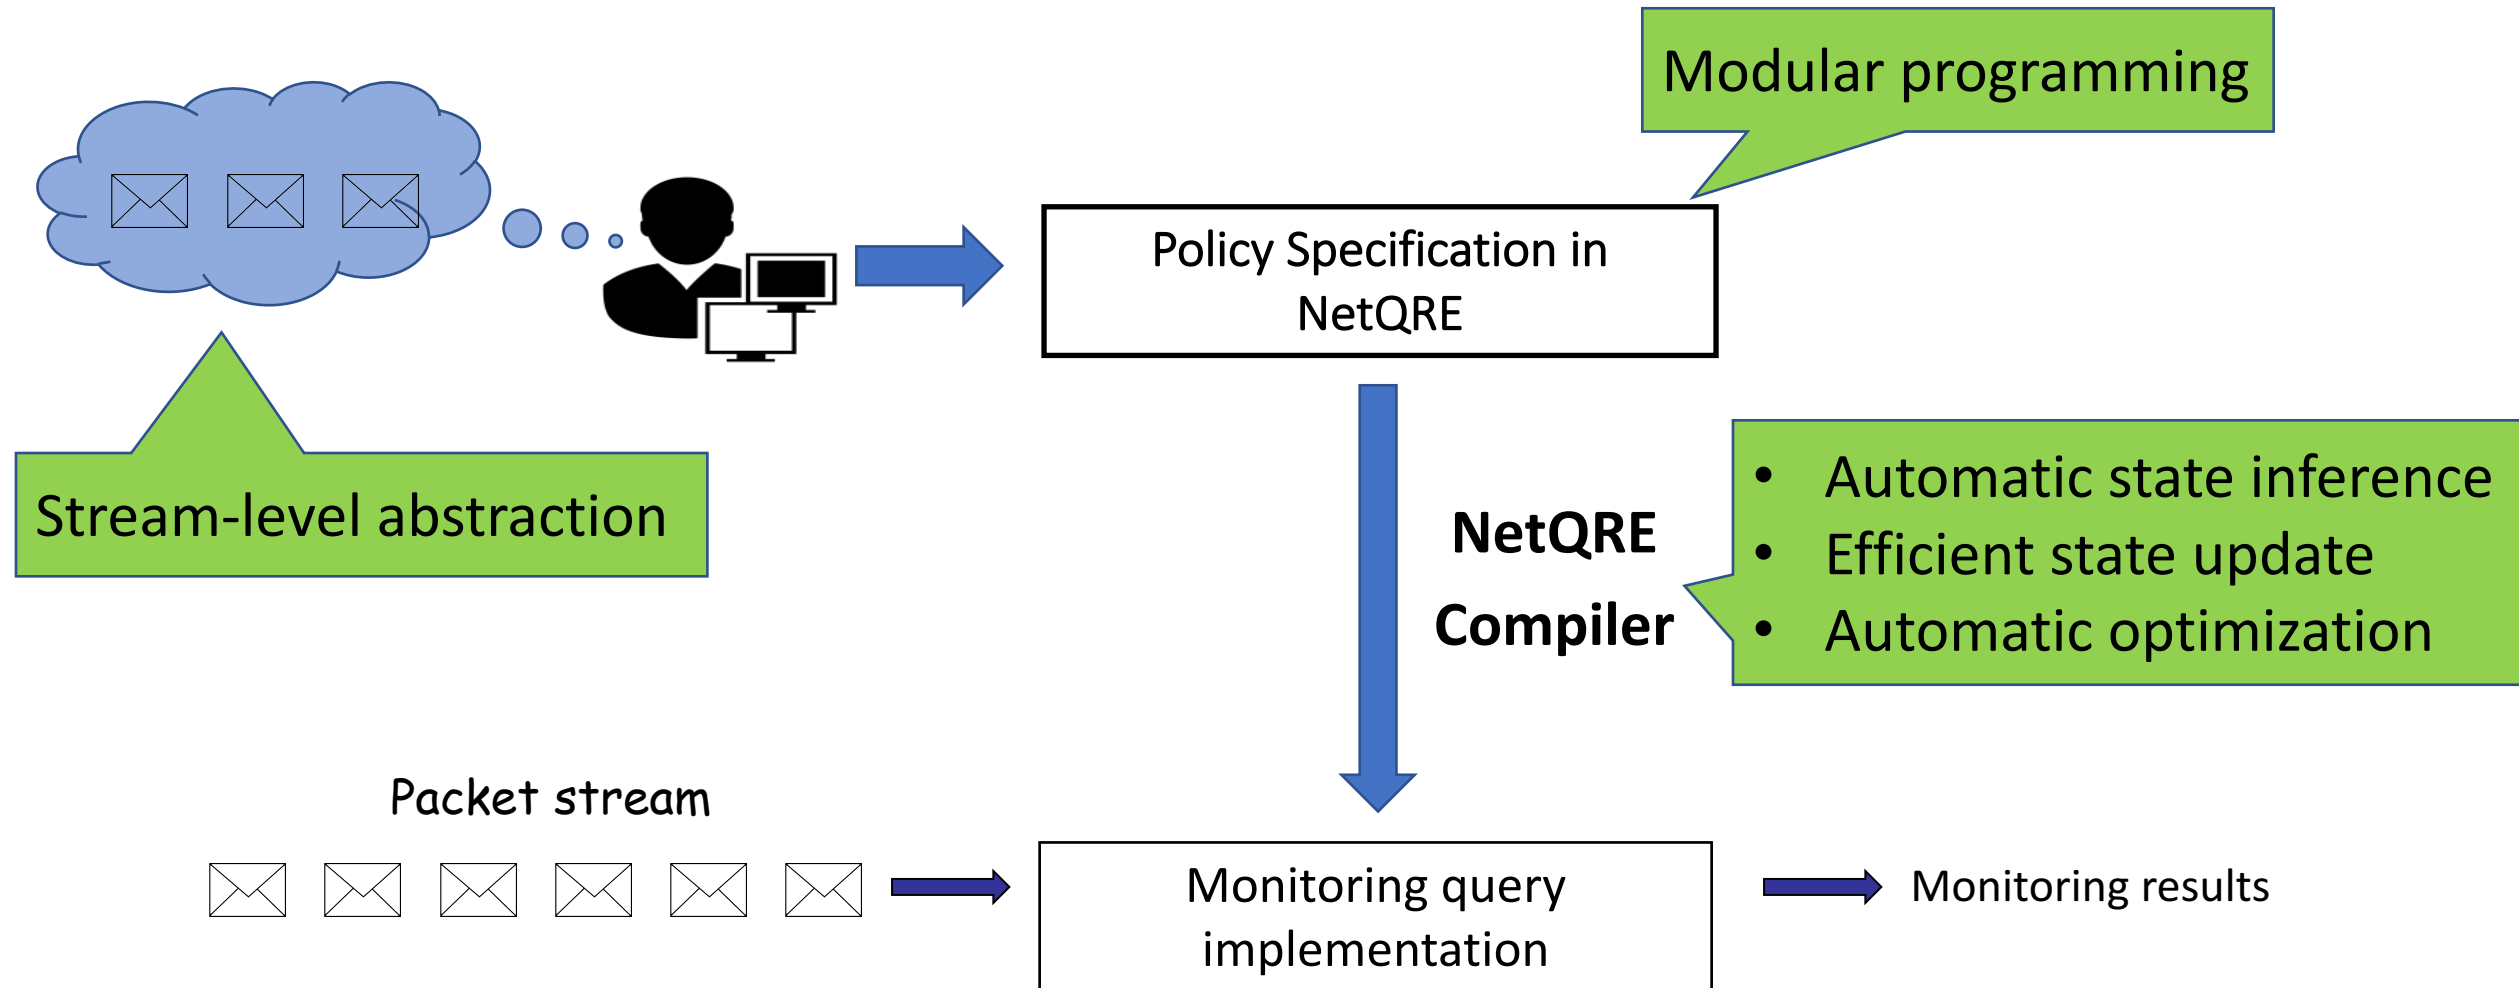
Stateful: Need to maintain state to track VoIP sessions with each incoming packet

Quantitative: Need to compute numerical aggregate based on metrics of past history and across users



What low-level state to maintain?  
How to update it?

# NetQRE Overview



# Outline

- Motivation
- NetQRE language
- NetQRE compiler
- Implementation
- Evaluation

# Outline

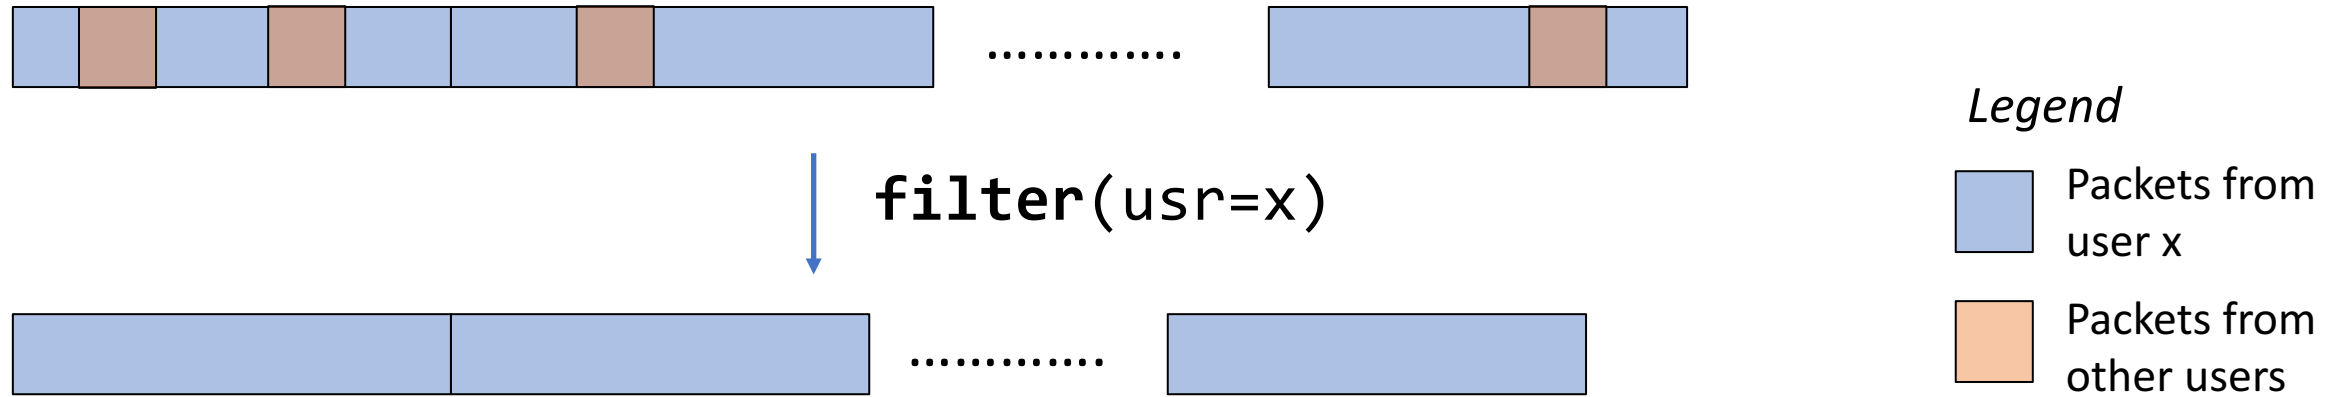
- Motivation
- NetQRE language
- NetQRE compiler
- Implementation
- Evaluation

# Modular Programming of VoIP Monitoring

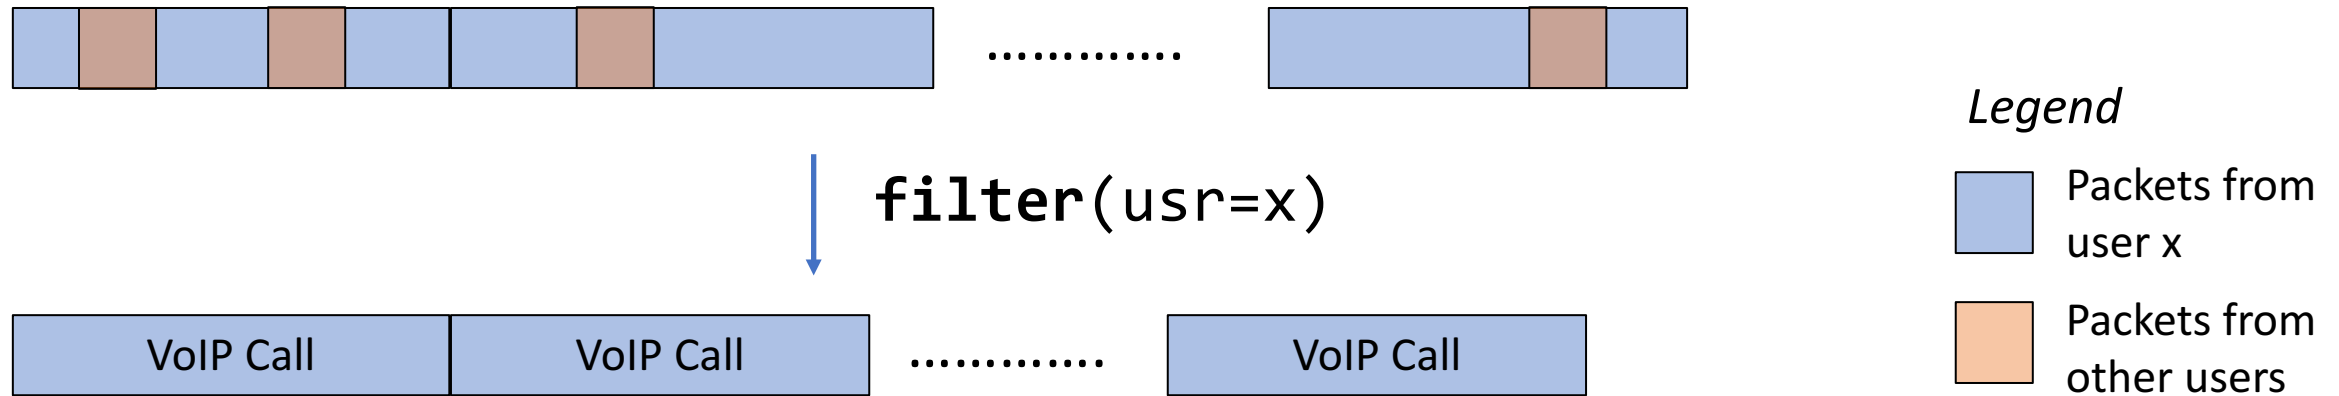
- Input: packet stream from all users
- Output: average number of VoIP calls per user
- Procedure:
  - Step 1: Focus on the packet stream from an arbitrary user  $x$
  - Step 2: View the stream as a sequence of calls, and identify each call
  - Step 3: Aggregate across all calls in the stream of the user
  - Step 4: Aggregate across all users



# Step 1: Focus on Packet Stream of User x



## Step 2: Identify A Call



How to specify the pattern of a VoIP call?


## Step 2: Identify A Call





↓  
**filter(usr=x)**



*Legend*

 Packets from user x

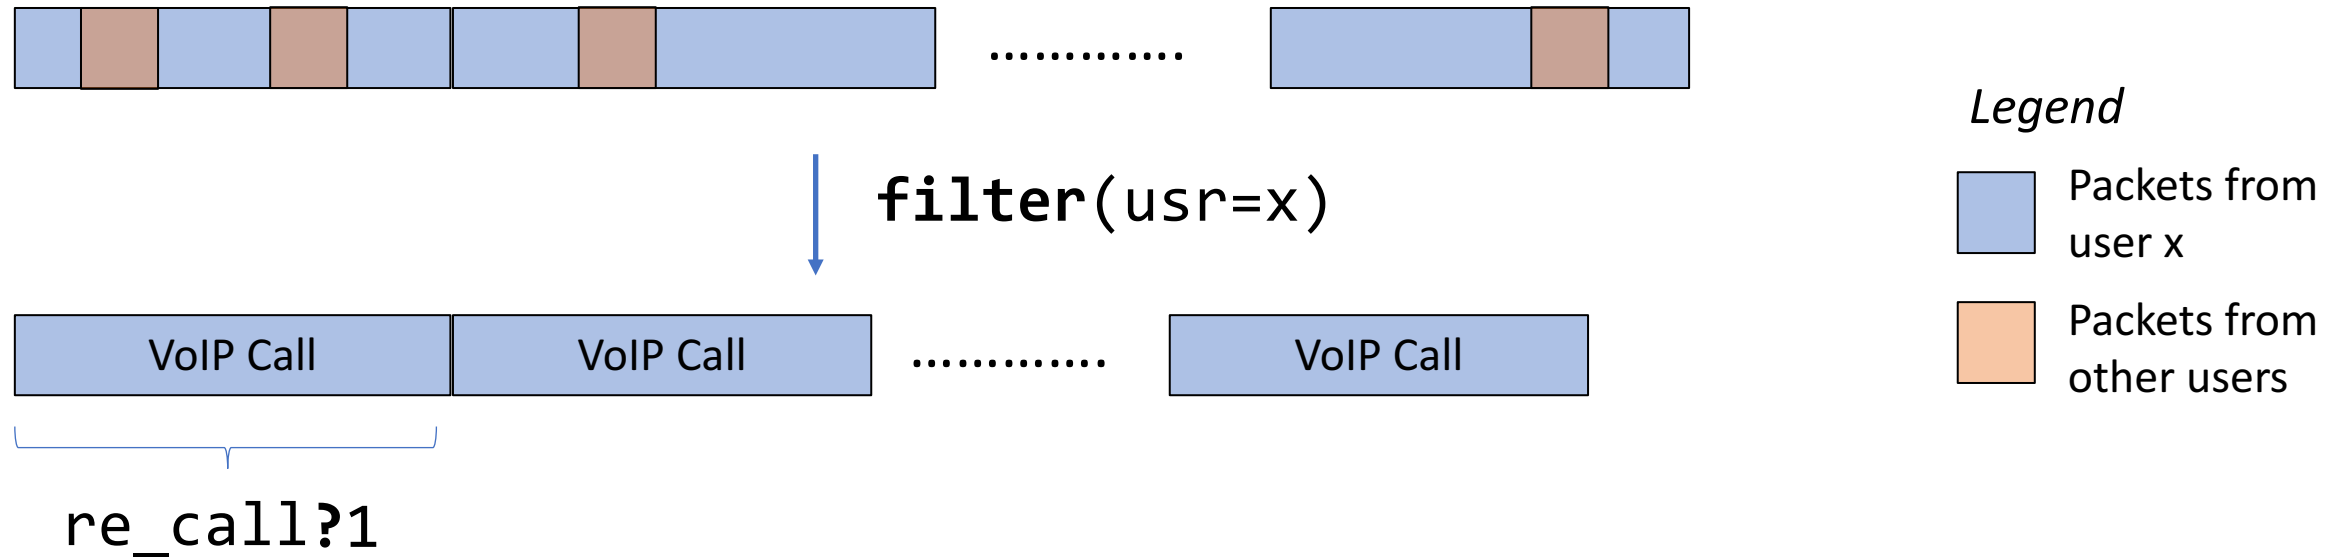
 Packets from other users



regular pattern

re\_call = [invite] [200] [ack] [data]\* [BYE] [200]

## Step 2: Identify A Call



How to associate a numerical value with each pattern?

## Step 2: Identify A Call



↓  
**filter(usr=x)**





`call = re_call?1`

`call`

`call`

*Legend*

 Packets from user x

 Packets from other users

# Step 3: Aggregate over All Calls



↓ **filter(usr=x)**




`call = re_call?1`


`call`

`call`

`iter(call, sum)`


*Legend*

 Packets from user x

 Packets from other users

# Stream Iteration: **iter**( $f$ , aggop)

- $f$  is a NetQRE function
- aggop is an aggregation operator, such as sum, avg, max, min.
- Split the stream into multiple substreams  $s_1, \dots, s_n$  such that  $f$  is defined on each  $s_i$
- Returns  $\text{aggop}(f(s_1), \dots, f(s_n))$



## Step 4: Aggregation over All Users



**filter**(usr=x)



call


call


call

`call_usr(x) = iter(call, sum)`

`avg{ call_usr(x) | User x }`

*Legend*

 Packets from user x

 Packets from other users



# Requirements & Key Ideas

Requirements	Key Ideas
Pattern matching for recognizing traffic patterns	Regular expression (RE) for pattern matching
Handle arbitrary & unknown value	Parametric extension to RE
Quantitative aggregations	Quantitative extension to RE

# NetQRE Language

## Regular Expression

- Atoms: letters
  - E.g. a, b, ...
- Base RE: atoms
- Union:  $f \mid g$
- Concatenation:  $f \circ g$
- Kleene star:  $f^*$

Parametric  
extension



Quantitative  
extension




## NetQRE

- Atoms: predicate over packets
  - E.g. [srcip==x], [dstip!=10.0.0.1]...
- Base NetQRE:  $re \ ? \ v$
- Choice:  $re \ ? \ f : g$
- Split: **split**(f, g, aggop)
- Iteration: **iter**(f, aggop)
- Aggregation over parameter:  
aggop{ f(x) | Type x }
- Streaming composition:  $f \gg g$

Details in the paper

# Stream Split: **split**( $f$ , $g$ , $\text{aggop}$ )

- $f$  and  $g$  are two NetQRE functions
- $\text{aggop}$  is an aggregation operator, such as `sum`, `avg`, `max`, `min`.
- Split the stream into two substreams  $s_1$  and  $s_2$ , such that  $f$  is defined on  $s_1$  and  $g$  is defined on  $s_2$
- Returns  $\text{aggop}(f(s_1), g(s_2))$



# Outline


- Motivation
- NetQRE language
- NetQRE compiler
- Implementation
- Evaluation

# NetQRE Compilation


- Goal: Evaluating a query online with small state
  - Independent of length of packet stream
- Insight: Leverage compilation of regular expression to DFA
- Question 1: How to handle parameters?
  - Insight: Lazy instantiation
- Question 2: How to evaluate `split(f,g,aggop)` and `iter(f,aggop)` online?
  - Insight: Keep all possible (but bounded number of) cases
  - Details in the paper

# Compilation of RE with Parameters


`exist_src(x)`





Eager  
Instantiation



⋮

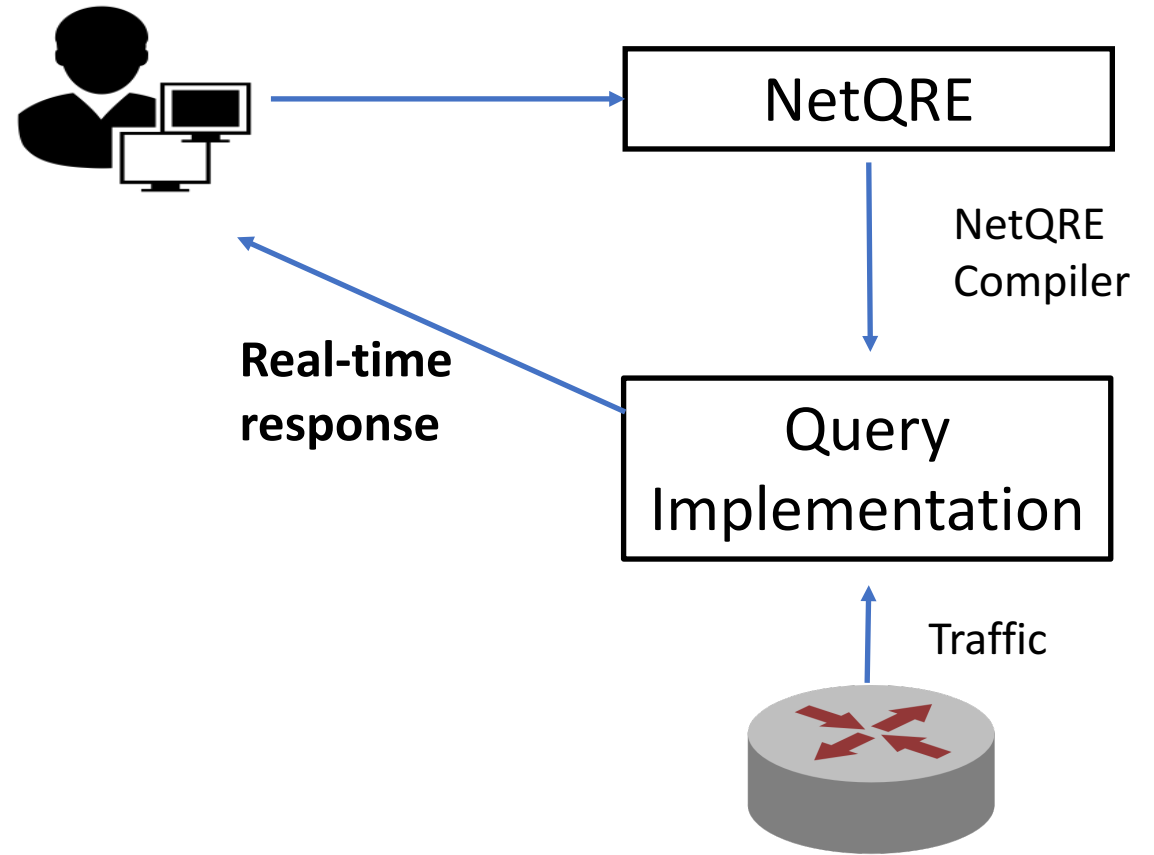


Lazy  
Instantiation



# Implementation

- Single-node deployment
- Compiler implemented in C++
- Compiled code C++
- pcap library for packet capturing



# Outline

- Motivation
- NetQRE language
- NetQRE compiler
- Implementation
- Evaluation



# Evaluation

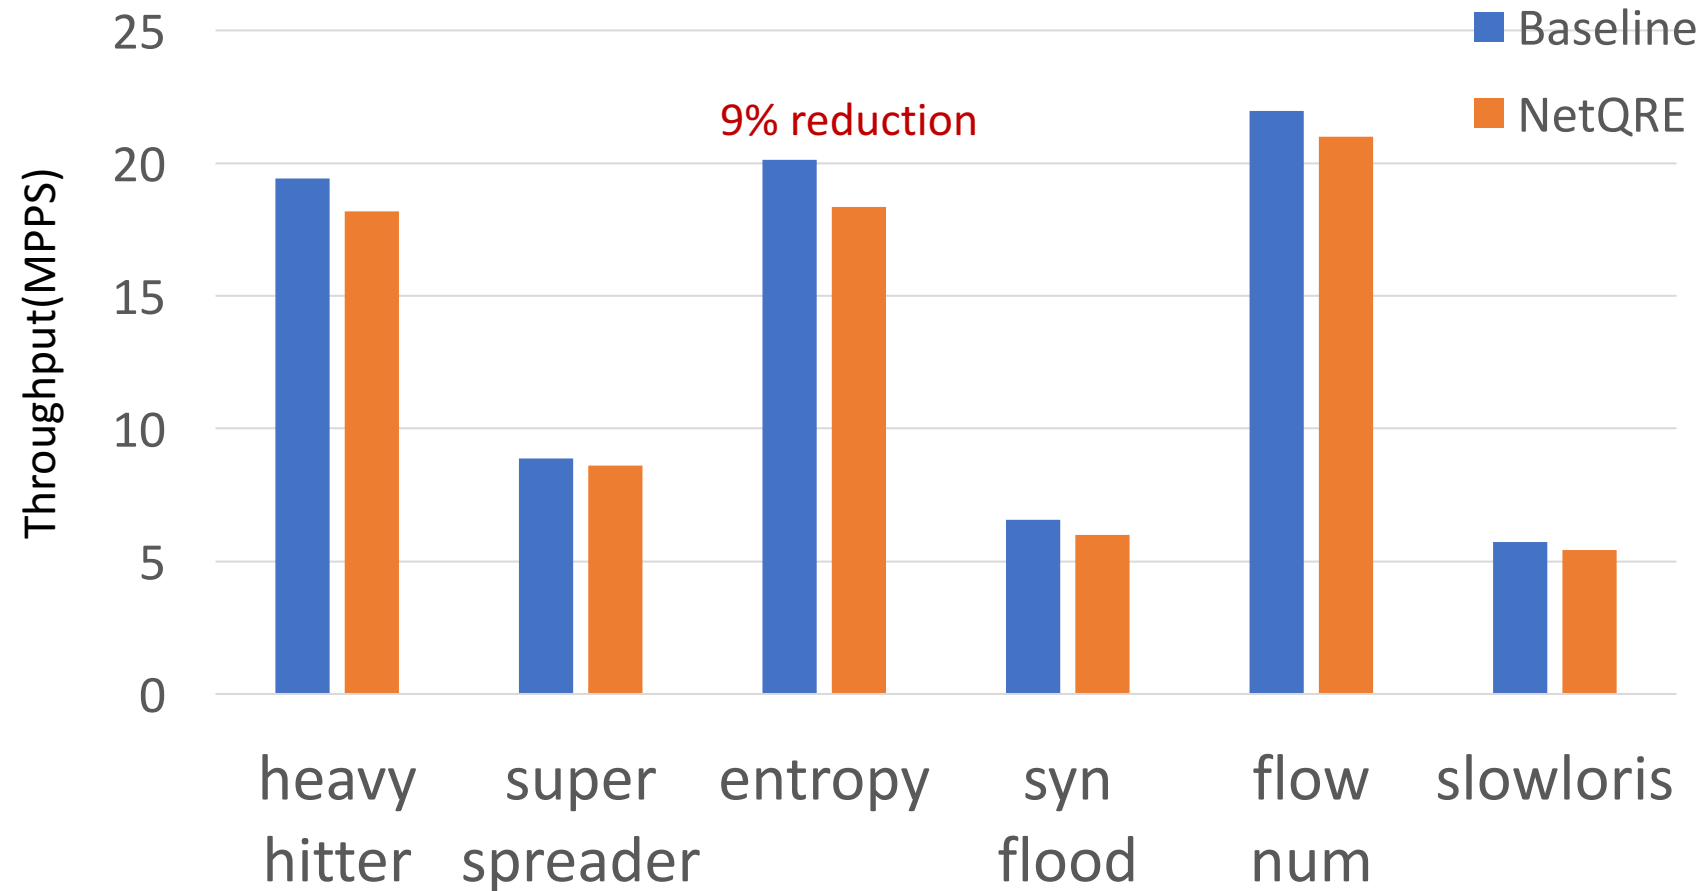
- Is the NetQRE language expressive?
- Is the NetQRE compiled implementation efficient?

# Evaluation: Expressiveness

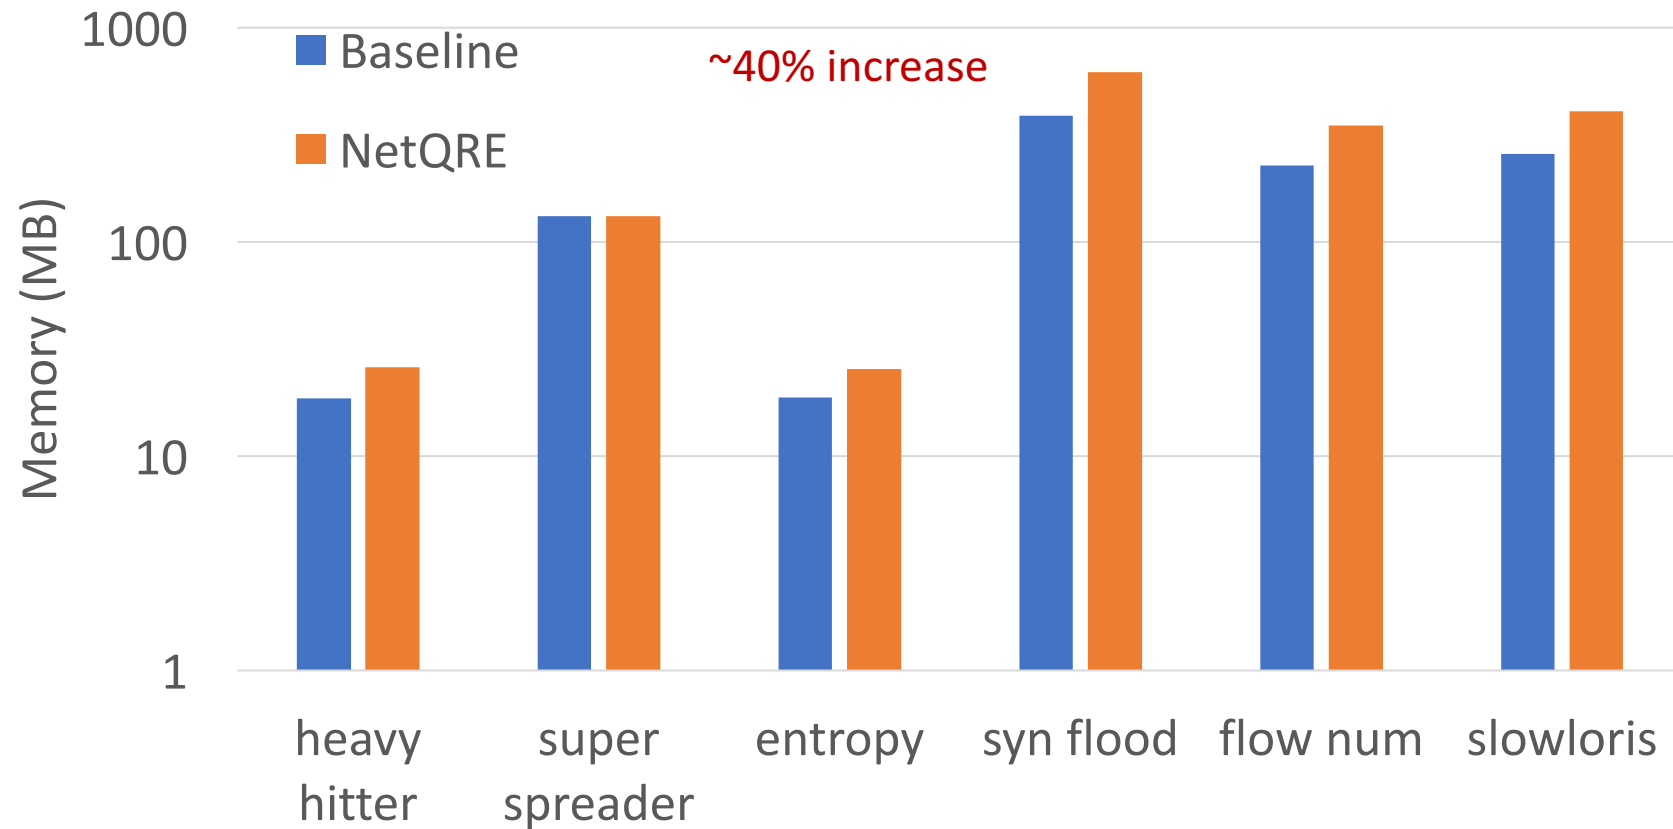
	LoC
Heavy hitter	6
Super spreader	2
Entropy estimation	6
Flow size distribution	8
Traffic change detection	10
Count traffic	2
Completed flows	6
SYN flood detection	9
Slowloris detection	12
Connection lifetime	8
Newly opened connection	11
# duplicated ACKs	5
# VoIP calls	7
VoIP usage	18
DNS tunnel detection	4
DNS amplification	4

- Expressive
- Concise
  - 100-1000+ LoC in compiled implementation
  - 100+ LoC in manual implementation

# Evaluation: Throughput



# Evaluation: Memory



# Conclusion

- Motivation: Network monitoring needs high-level abstractions
- Contributions:
  - Stream-level programming abstraction
  - Parametric and quantitative extension to regular expressions
  - Expressive to capture a wide range of monitoring policies
  - Compiled implementation efficient in both throughput and memory
- Future work:
  - Hardware implementation
  - Distributed deployment of NetQRE programs