



Le génie pour l'industrie

ACM SIGCOMM 2019 Workshop on Networking for Emerging
Applications and Technologies (NEAT 2019)

FlexNGIA

A Fully Flexible Novel Architecture for the Next-Generation Tactile Internet

Mohamed Faten Zhani

École de technologie supérieure (ÉTS Montreal)

University of Quebec

Canada

Beijing, China, 19 August 2019

Outline

- A Glance into the Future
- Limitations of Today's Internet
- FlexNGIA: Fully-Flexible Next-Generation Internet Architecture
- Use cases
- Conclusion

M. F. Zhani, H. ElBakoury, "*FlexNGIA: A Flexible Internet Architecture for the Next-Generation Tactile Internet*," ArXiv 1905.07137, May 17, 2019
<https://arxiv.org/abs/1905.07137>

A Glance into the Future

Future Applications

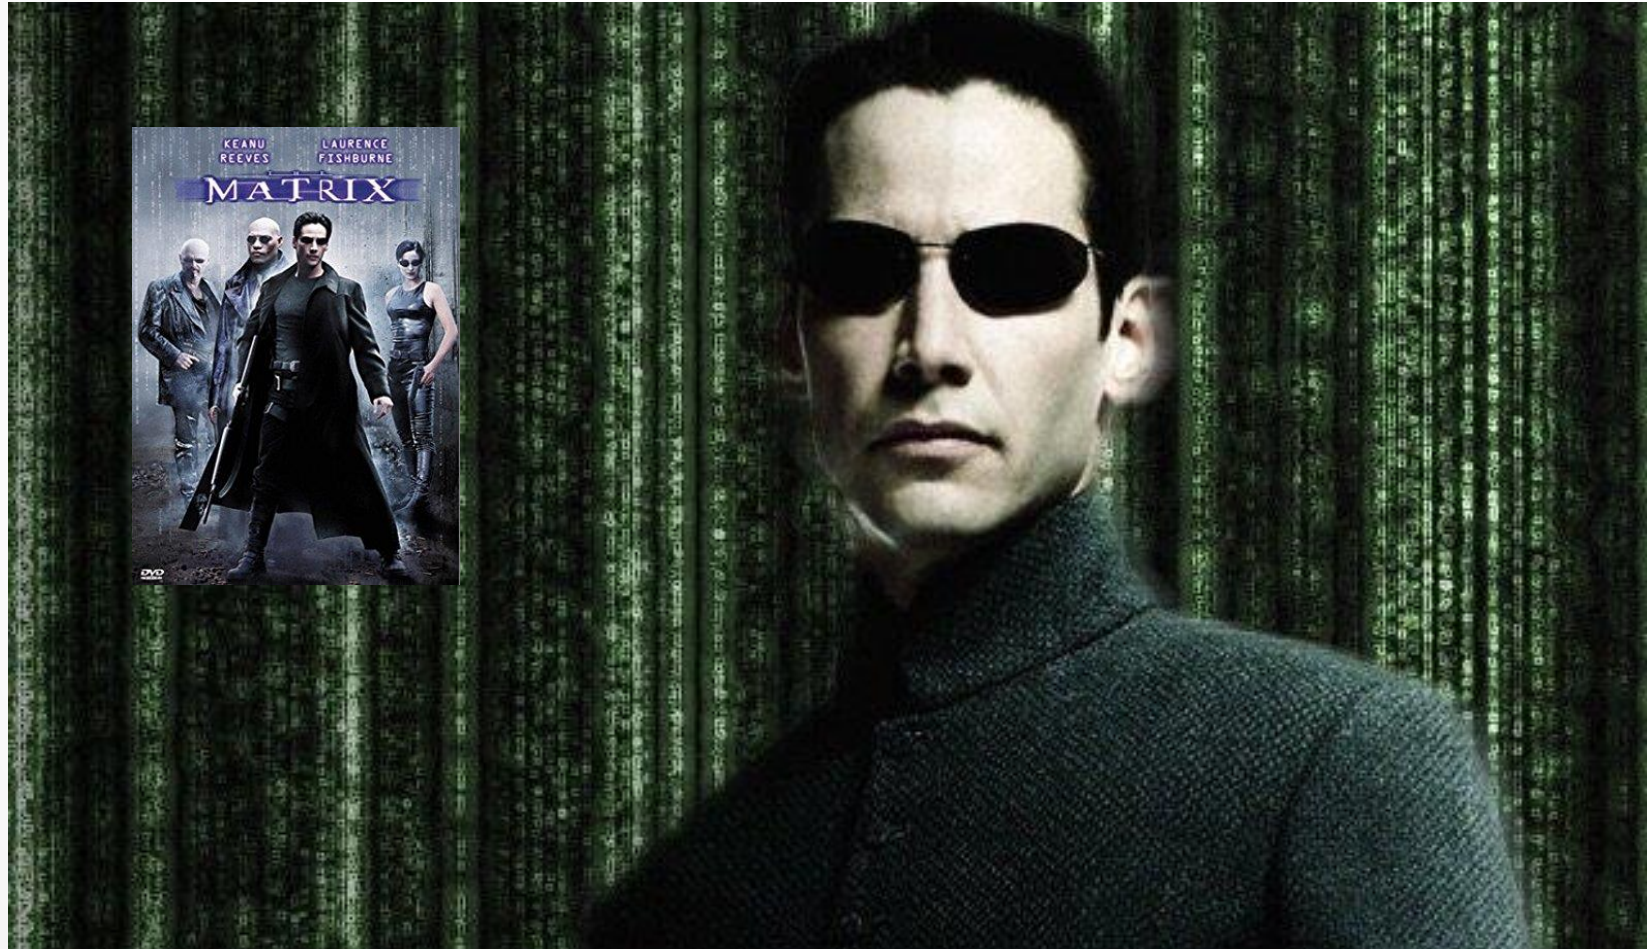
- Telepresence
- Virtual Reality
- Augmented Reality
- Holoportation
- Haptics
- ...



Loading the Matrix...



Welcome to the Matrix



Future Applications Requirements & Characteristics

- Characteristics
 - Octopus-like applications: huge number of flows for each application
 - Changing requirements : requirements can change over time
- Requirements:
 - High processing power: real-time processing
 - High bandwidth (e.g., VR (16K, 240 fps) → 31.85 Gbps)
 - Ultra-low Latency: 1ms to 20ms
 - Multi-flow synchronization
 - High availability

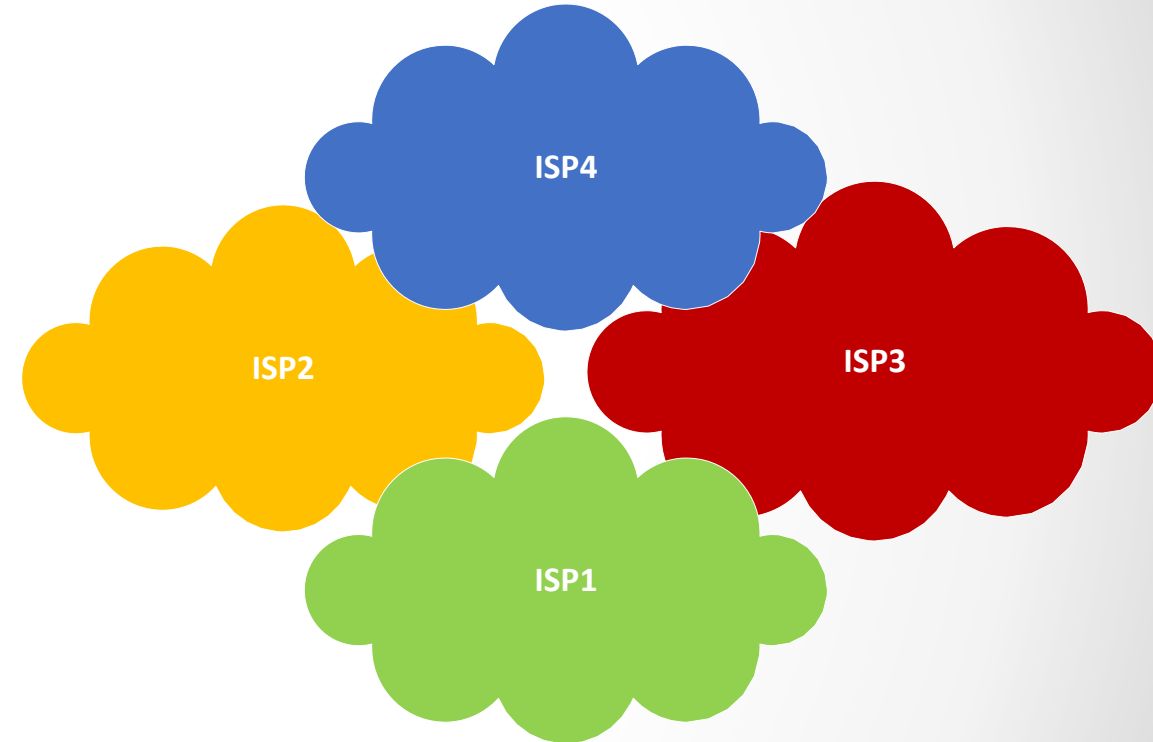


Outline

- A Glance into the Future
- Limitations of Today's Internet
 - Internet Infrastructure and Services
 - Network Stack Layers and Headers
- FlexNGIA: Fully-Flexible Next-Generation Internet Architecture
- Use cases
- Conclusion

Internet Infrastructure and Services

- A network of networks
- Offered service: “Best effort” data delivery.. no more
- No control over the infrastructure
 - ➔ No control over the end-to-end path and quality of service
 - ➔ No performance guarantees



Transport Layer Protocols

Many modern protocols like SCTP and QUIC but let's focus first on TCP..

- One-size-fits-all service offering: TCP offers reliability, data retransmission, congestion and flow control
- Blind Congestion control
- The two end points limitation:
 - High retransmission delays ($\sim 3x$ e2e delay)
 - Transport and network layers are not aware which flows belong to the same application

Network Layer Protocols

- Not aware of the applications
 - The application composition (in terms of flows)
 - Performance requirements of each of these flows and how these requirement change over time
 - ➔ Drop packets « blindly »
- No collaboration with the transport layer
 - Do not provide explicit feedback or support to transport layer (maybe ECN is interesting but it is not enough)
 - Do not help with other transport services (e.g., reliability)

Network Stack header

Problems with current headers:

- Do not provide additional informations about objects/sensors, flows belonging to the same application, applications' requirements, etc.
- Not flexible enough: It is not easy to incorporate meta-data and commands

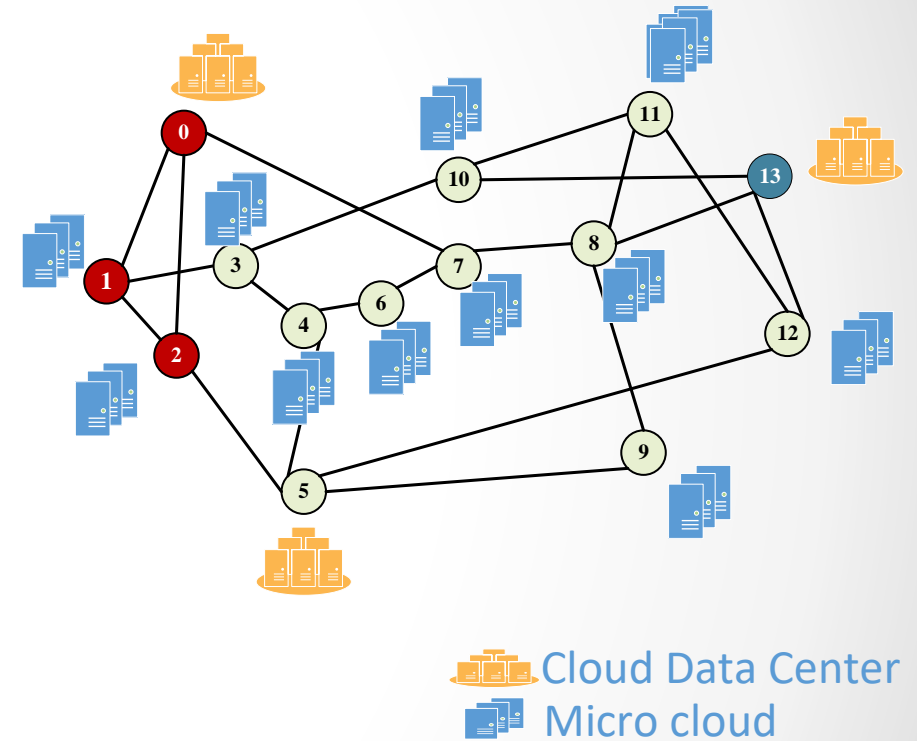
Outline

- A Glance into the Future
- Limitations of Today's Internet
- FlexNGIA: Fully-Flexible Next-Generation Internet Architecture
 - Future Internet Infrastructure and services
 - Business Model
 - Management Framework
 - Network Protocol Stack/Functions
 - Stack Headers
- Use cases
- Conclusion

Future Internet Infrastructure and Services

How a network will look like?

- Computing resources are everywhere: Available at the edge and at the core of the network
- Commodity servers but also dedicated hardware, FPGA, GPU, NPU, etc.
 - ➔ In-Network computing
 - ➔ Reduce steering delay
 - ➔ Full Programmability: Any function could be provisioned anywhere (virtual machines/containers)



Future Internet Infrastructure and Services

How does Future Internet look like?

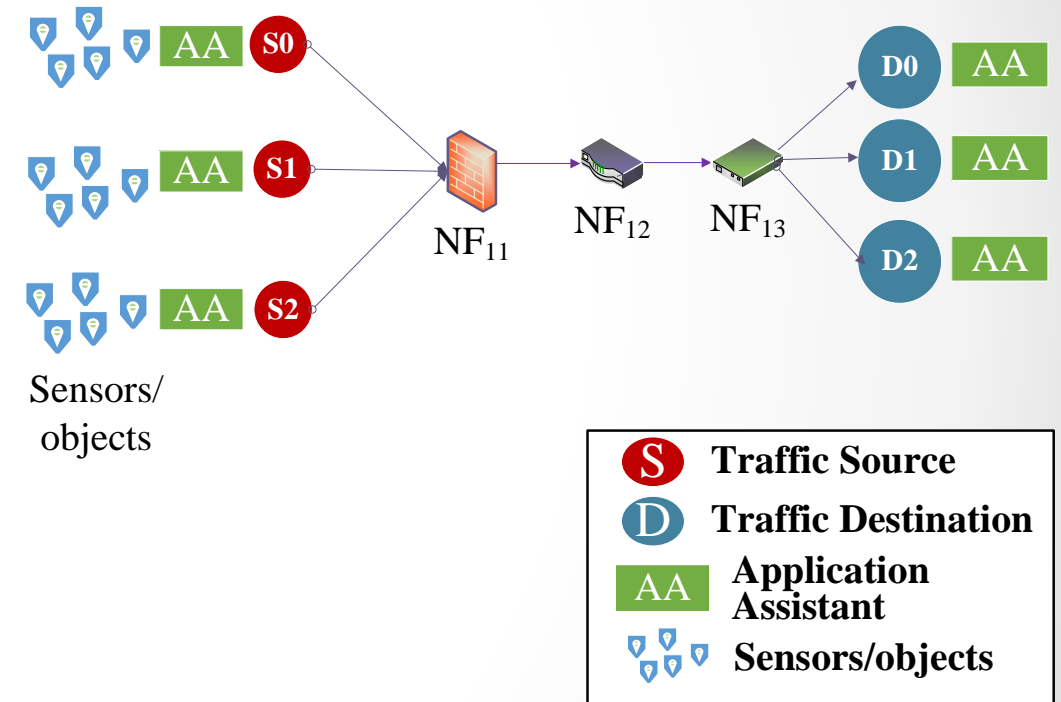
- Still a network of networks..
- What is new?
 - More services: Service Function chains
 - ➔ More advanced functions
 - ➔ More than just delivery
 - Stringent performance guarantees



Future Internet Infrastructure and Services

Service Function Chain (SFC)

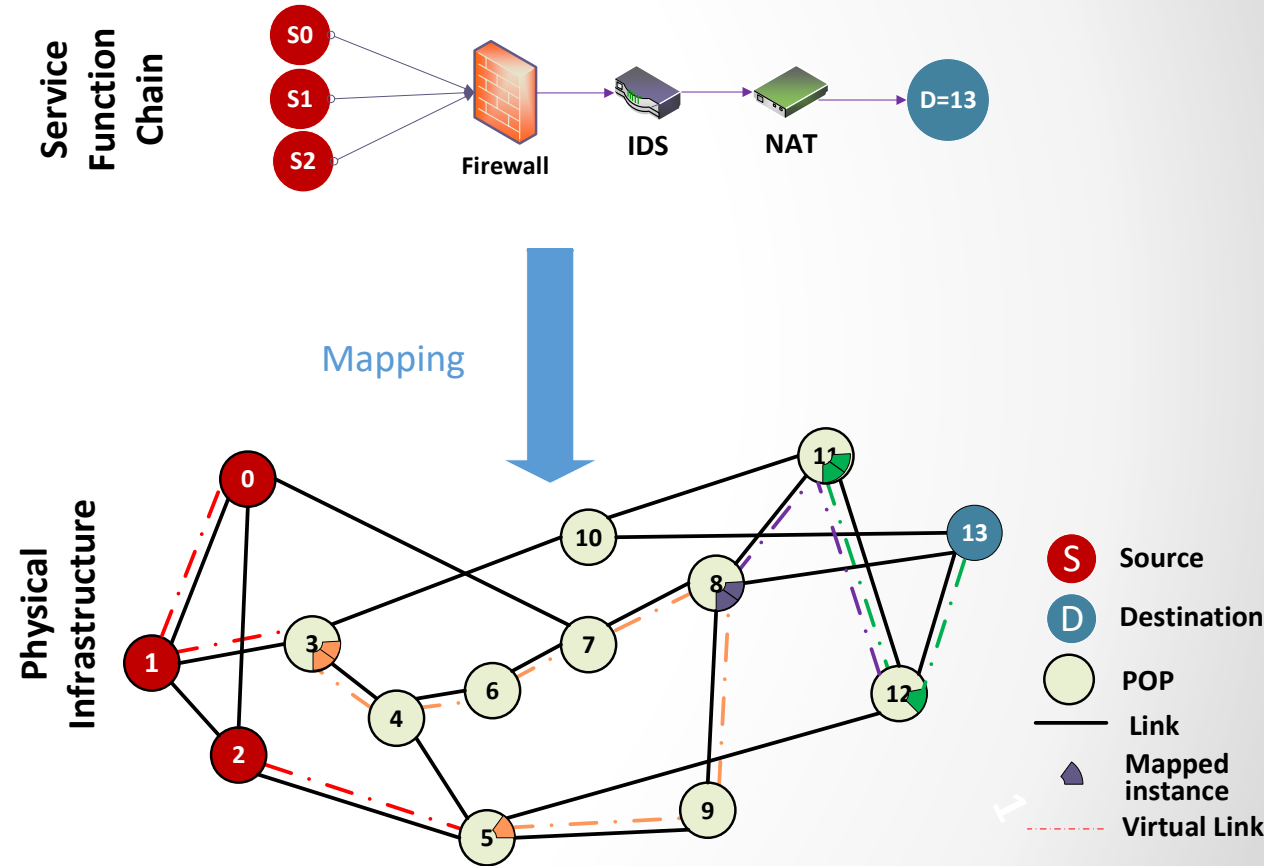
- Multiple connected network functions
- Multiple sources and destinations
- Made out from Network Functions
- Defines, for each network function, the type, software, input/output packet format, expected processing delay, buffer size
- Defines performance requirements (e.g., throughput, packet loss, end-to-end delay, jitter)



Business Model

Network Operators

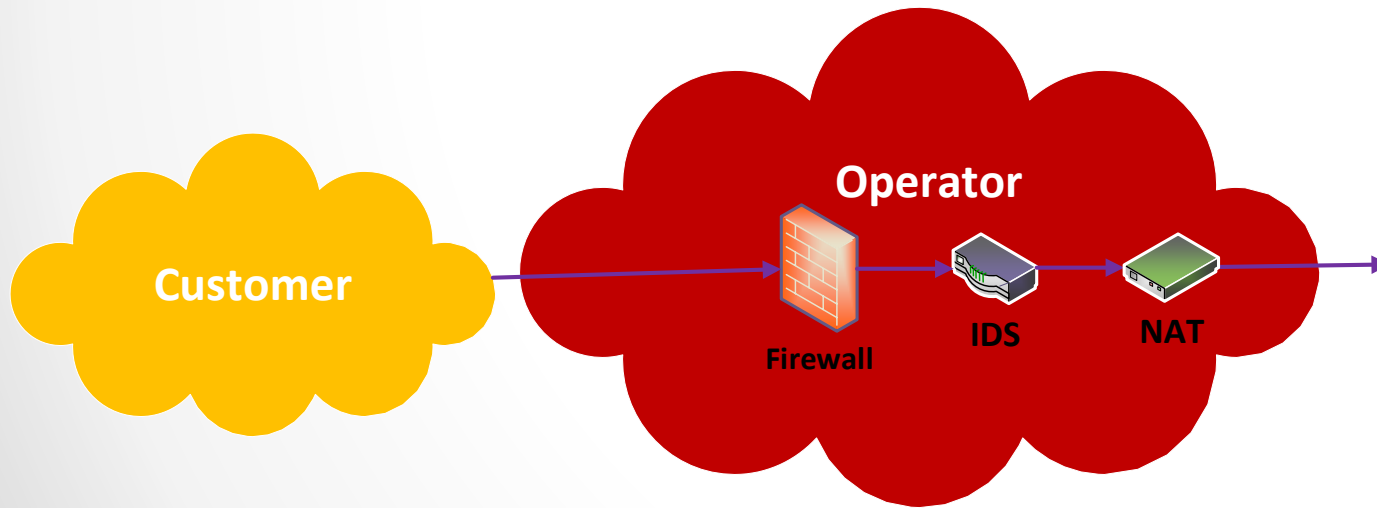
- Own and manage the physical infrastructure (i.e., one network)
- Deploy platforms and software required to run network functions
- The service could be simply data delivery or a SFC
- Provision and manage SFCs



Business Model (cont)

Customers

- Could be other network operators, companies or Institutions
- Define the required SFC and Identify the chain sources/destinations
- Rely on the operator to provision and manage the SFC and satisfy SLA

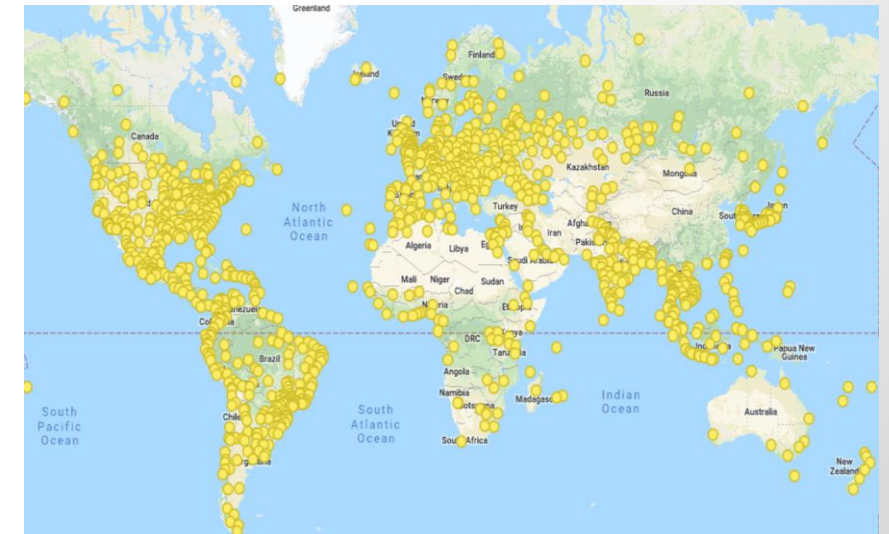
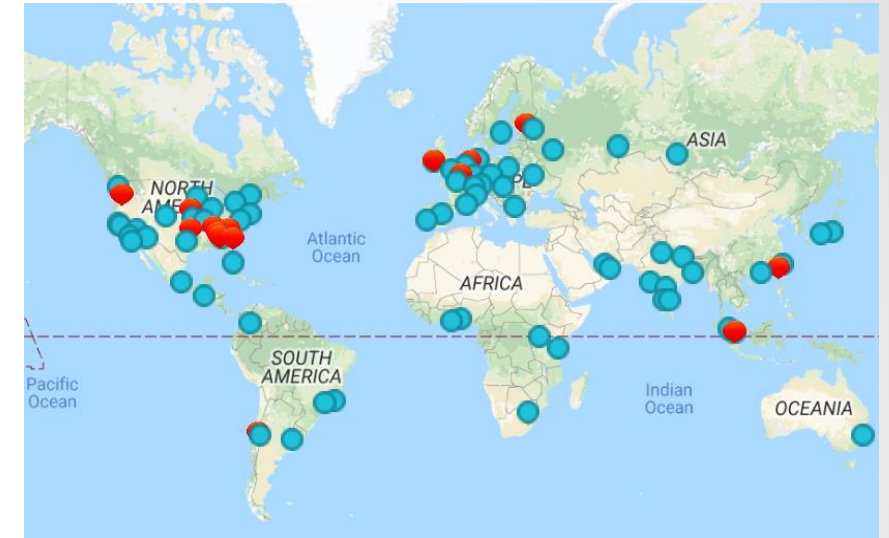
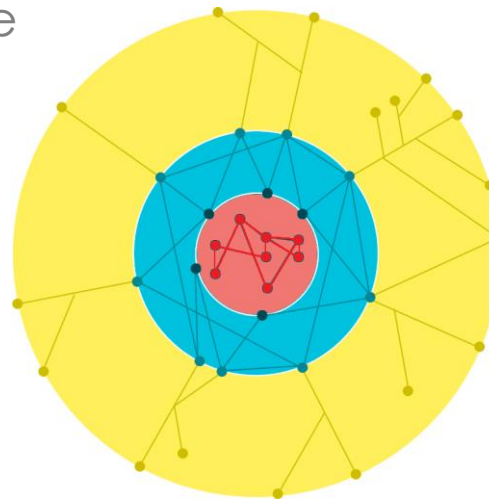


- SFC composition
- SLA requirements for the SFC
 - Bandwidth
 - End-to-end delay
 - Reliability, availability
- SLA requirements for each NFs
 - Processing power
 - Packet format(s)
 - Packet drop criteria...

Business Model (cont)

- Example of potential Network Operators:
 - ISPs (e.g., AT&T or Bell Canada) and web-scale companies (e.g., Google, Facebook, Amazon)
 - Example: Google Cloud Platform
 - World wide global Infrastructure
 - Software defined platform
 - Full control over the infrastructure

- 15 Data centers
- 100 Points of Presence (PoPs)
- 1000+ Edge nodes

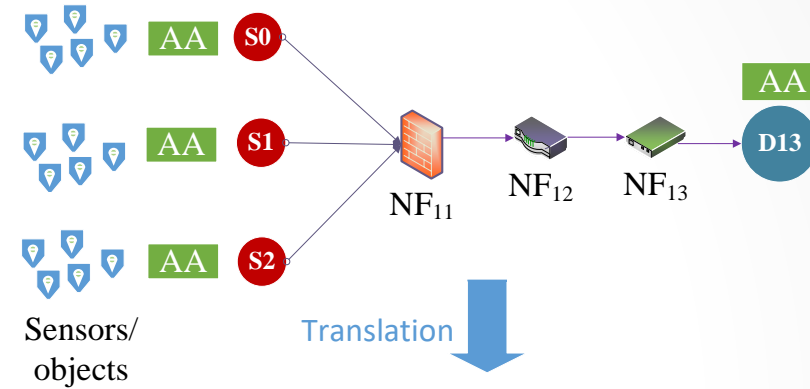


Resource Management Framework

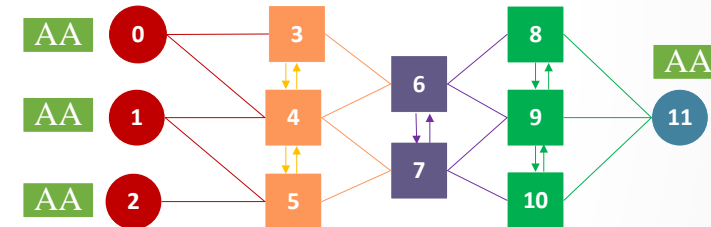
Resource Allocation

- The Service Function Chain (SFC) is defined by the application designer
- 2-step resource allocation:
 - Translation: the SFC is translated into a virtual topology
 - Mapping: virtual topology are mapped

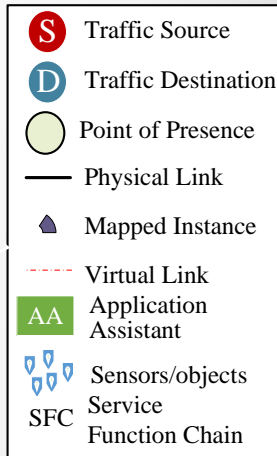
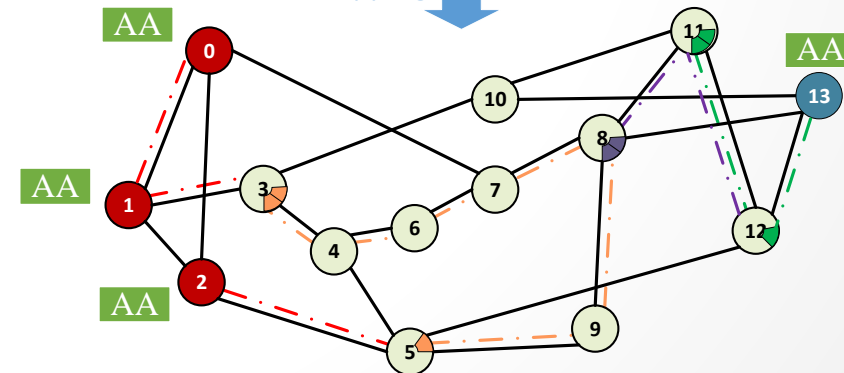
Service Function Chain SFC₁
associated with Application 1



Virtual
Topology



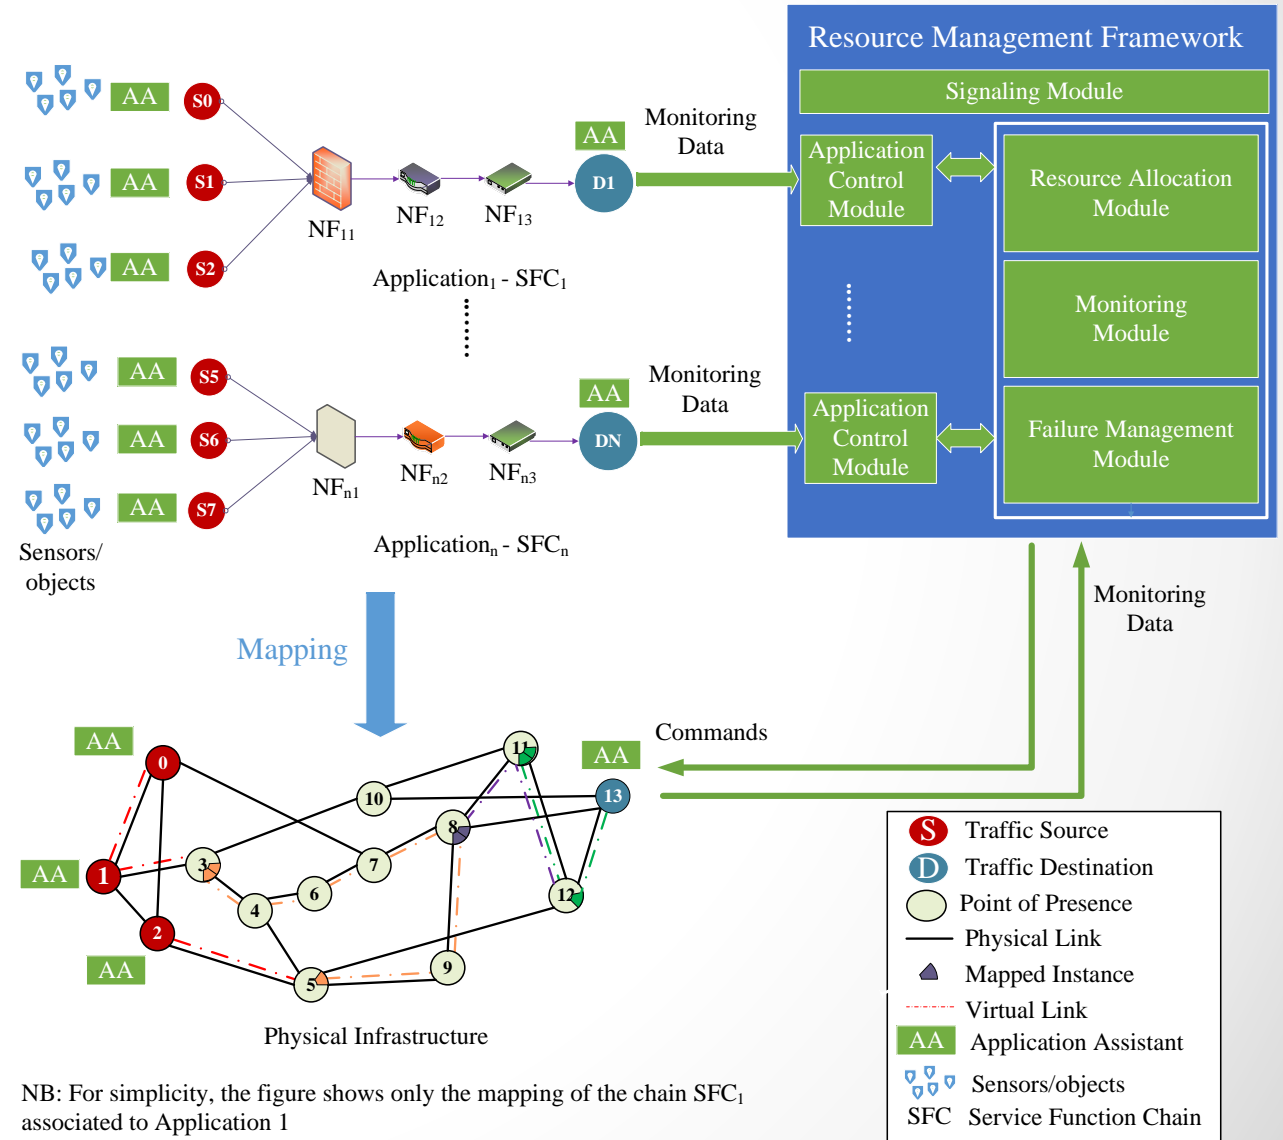
Physical
Infrastructure



Resource Management Framework

Main components:

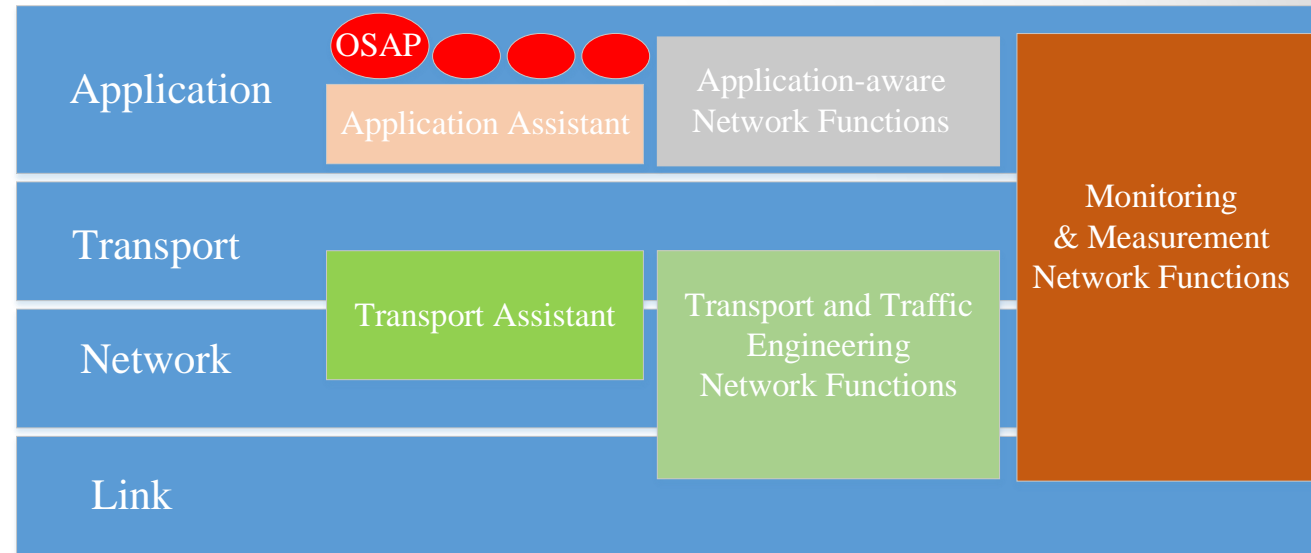
- Signaling module
- Application Control Module
- Ressource allocation Module



NB: For simplicity, the figure shows only the mapping of the chain SFC₁ associated to Application 1

Network Protocol Stack/Functions

- Basic Network Functions (e.g., packet forwarding)
 - Advanced Network Functions:
 - Could operate at any layer
 - Only limited by our imagination
 - Examples: packet grouping, caching and retransmission, data processing (e.g., image/video cropping, compression, rendering, ML), application-aware flow multiplexing (e.g., incorporating/merging data)
- ➔ Functions could break the end-to-end principle
- ➔ SDN++: SDN should go beyond configuring forwarding rules and should provide the ability to dynamically configure these new functions

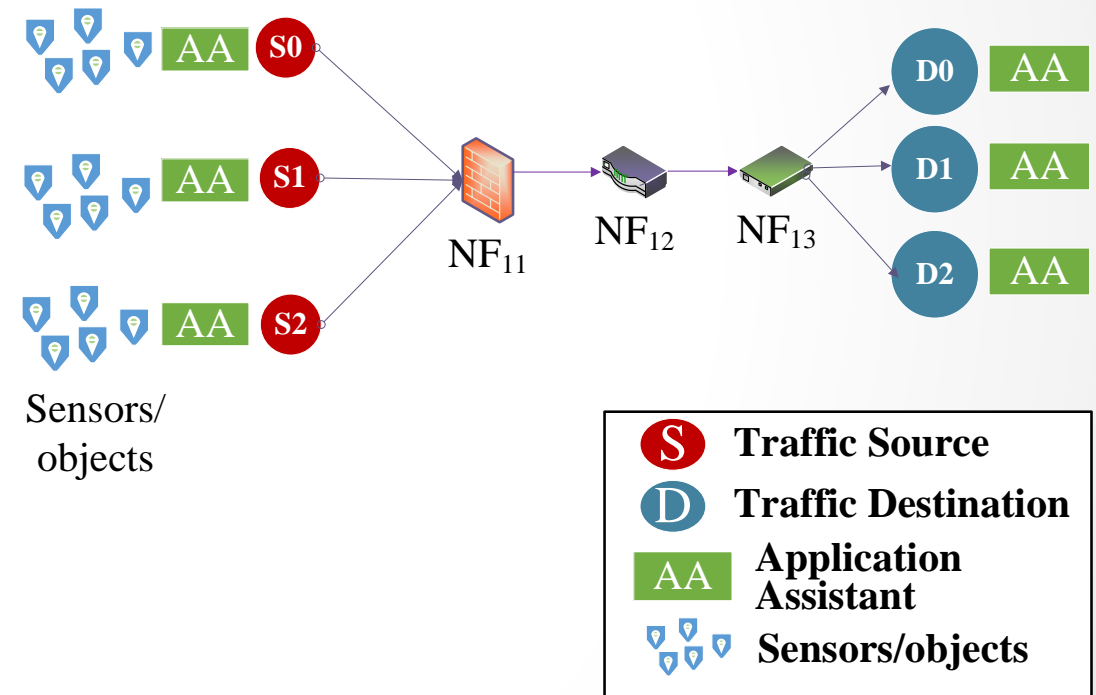


Network Protocol Stack/Functions

Application Assistant

Application Assistant (AA)

- One AA at each end-point
 - Interfaces with objects/sensors
 - Measures the application performance and user QoE
 - Identifies the applications' requirements at run-time
 - Adds additional metadata To be used by subsequent Network Functions
- ➔ Application-Aware Network Services

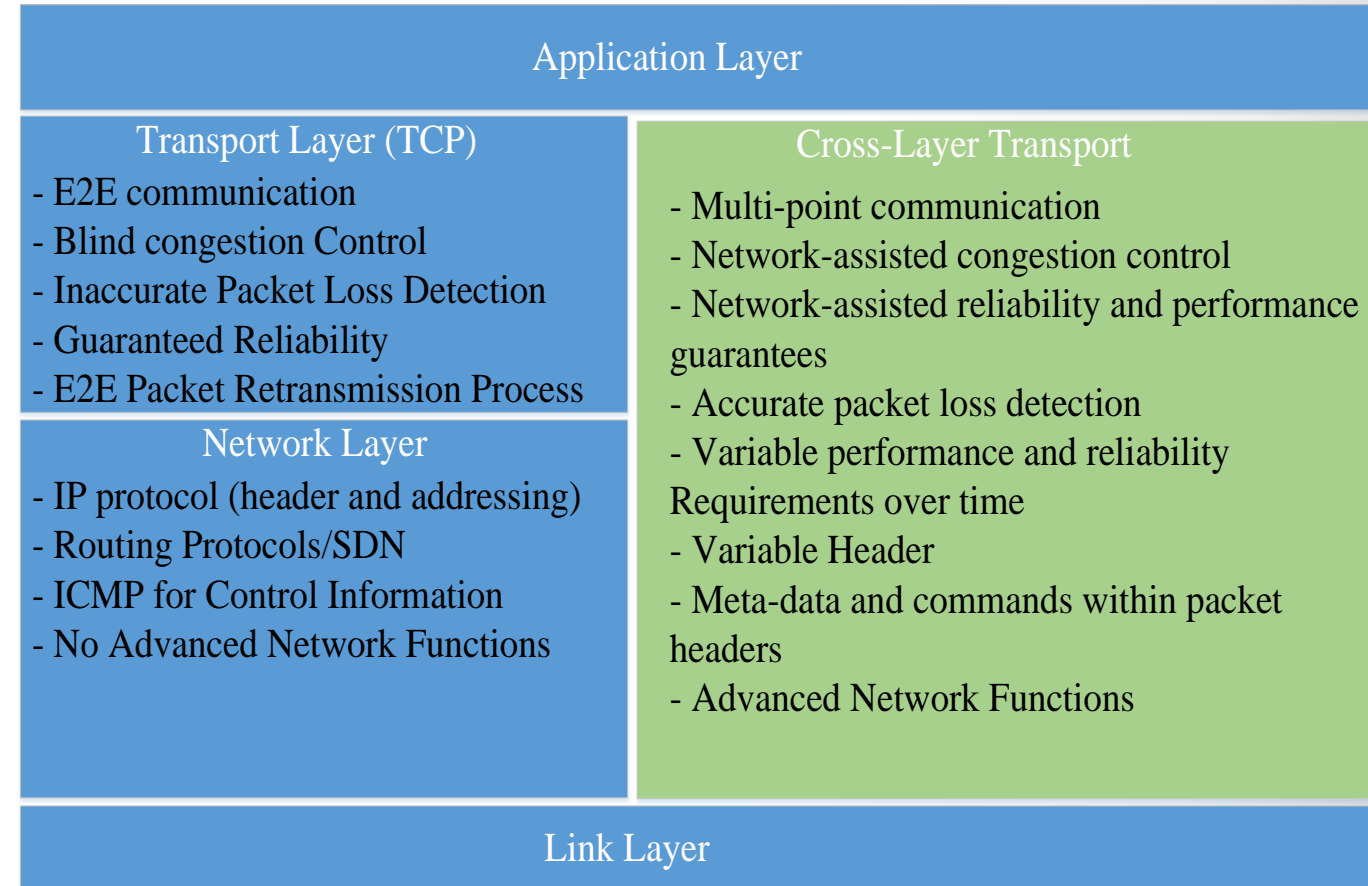


Network Protocol Stack/Functions

Transport Assistant

Transport Assistant (TA)

- A cross-layer Network Function
- Combines services of the transport and network layers
- Manages all the flows of the same application
- Implements Transport/Network functions (e.g., congestion control, packet loss detection, packet cache and retransmission, routing)
- One or multiple TA could be provisioned in the same SFC



Network Stack Headers

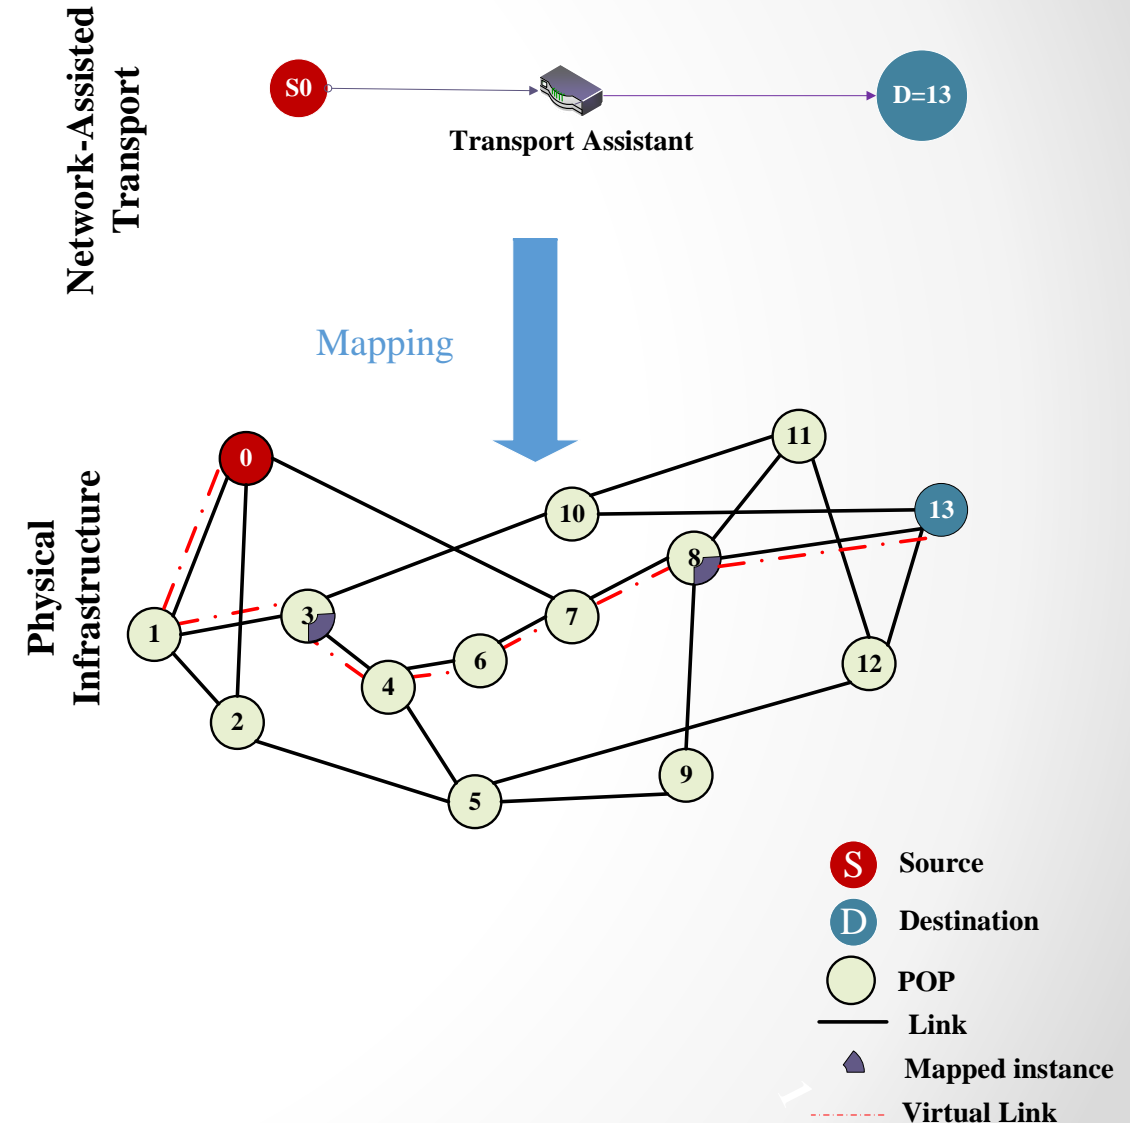
- Signaling packets
 - Instantiate an application
 - Convey application requirements
- Data packets: carry data
 - Layer 2 header: contains mainly the application id used for packet forwarding (similar to VLANs)
 - Upper layers:
 - Fully flexible header format (customizable meta-data and commands)
 - Defined depending on the application
 - Network functions should be aware of the expected format

Outline

- A Glance into the Future
- Limitations of Today's Internet
- FlexNGIA: Fully-Flexible Next-Generation Internet Architecture
- Use Cases scenarios
- Conclusion

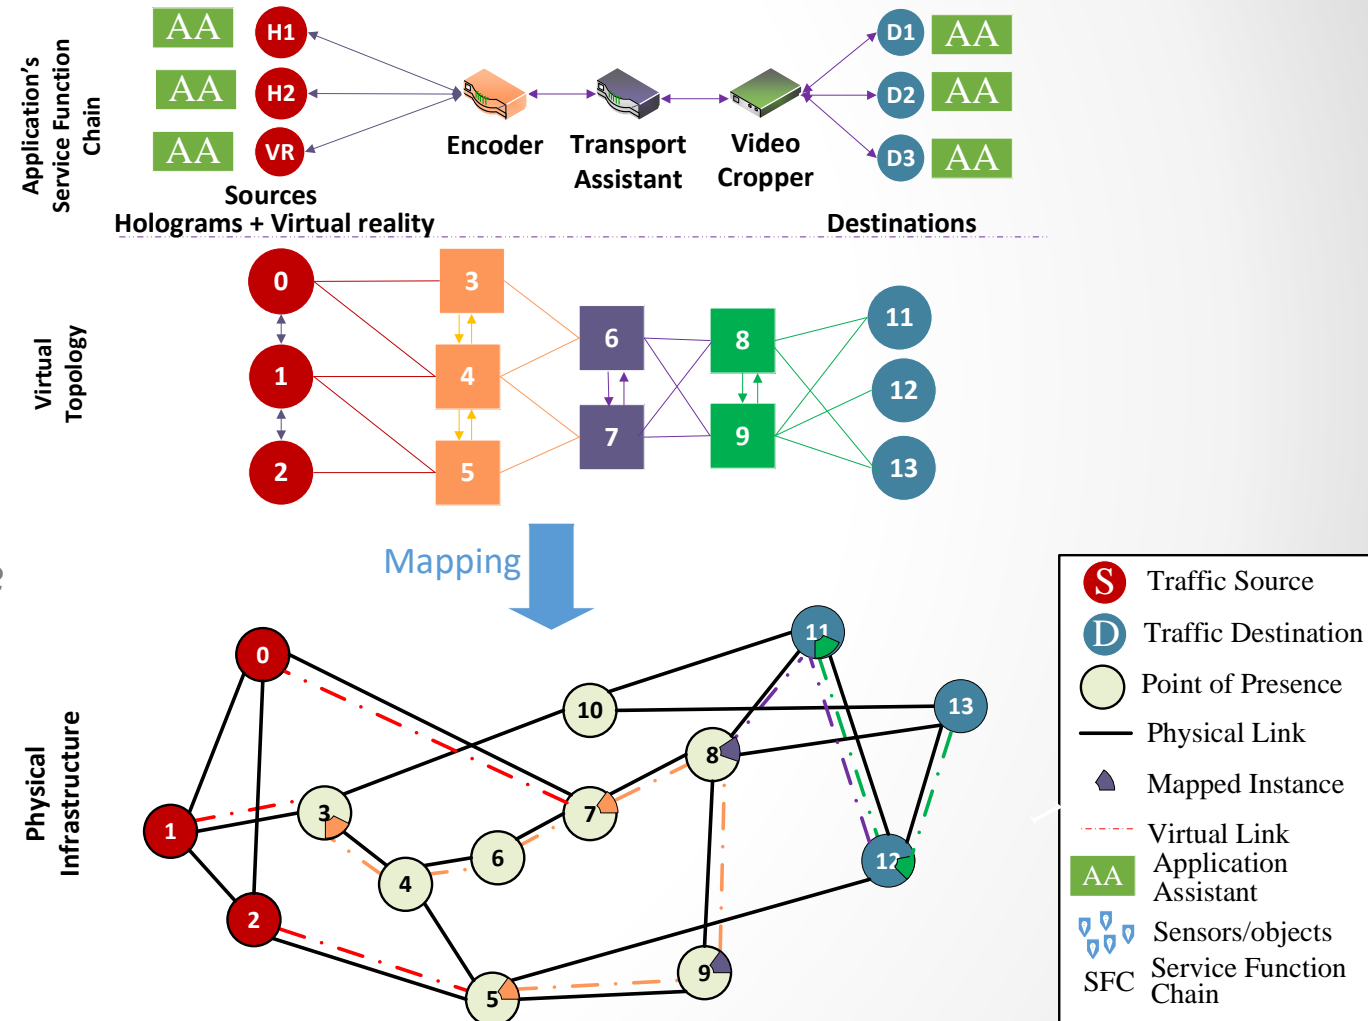
Network-Assisted Data Transport

- Goal
 - Minimize retransmission delay
 - Improved congestion control
- Solution: service chain with a "transport Assistant" function
- Service of the Transport Assistant:
 - Caching and retransmitting packets
 - Detecting packet loss
 - Congestion control: adjusting rate, dropping packets, compression



Mixed Virtual Reality and Holograms

- Users are exploring a virtual reality environment with several human holograms and objects
- Challenges
 - How many intermediate functions?
 - What kind of functions?
 - How the traffic should be steered from the flow sources?
 - How many instances for each function?
 - Where to place them?
- Example of deployment
 - Encoder: encode and compress video
 - Transport manager: congestion control
 - Video cropper: crop 3D objects



Outline

- A Glance into the Future
- Limitations of Today's Internet
- FlexNGIA: Fully-Flexible Next-Generation Internet Architecture
- Use Cases
- Conclusion

Conclusion

FlexNGIA

Computing
resources

Business model

Cross-layer Design
(Transport+Network)

Application-Aware
Network
Management

Flexible
headers

- In-Network
Computing:
any function
anywhere

- Multiple source
destination
Service Function
Chains
- Stringent
performance
requirements

- Breaking the
end-to-end paradigm
- In-network advanced
transport functions
- Better congestion control
- Stringent performance
and reliability guarantees

- Advanced functions
tailored to
applications
- App-aware traffic
engineering

- Tailored
to the
application

Looking for More Details?

- Mohamed Faten Zhani, <https://profs.etsmtl.ca/mfzhani/>
- FlexNGIA Project: M. F. Zhani, H. ElBakoury, "*FlexNGIA: A Flexible Internet Architecture for the Next-Generation Tactile Internet*," ArXiv 1905.07137, May 17, 2019 <https://arxiv.org/abs/1905.07137>



Thank You

Questions



Research Challenges

- Designing Service Function Chains tailored to applications
- High-performance softwarized functions
- Signaling
- Resource Allocation
- Fault-tolerance and Failure Management
- High-Precision and Fine-grained Monitoring and Measurements
- SDN++
- Distributed Cross-Layer Transport Protocol (sockets, caching, communication)
- Security and Privacy

Details are available in the paper (<https://arxiv.org/abs/1905.07137>)

Transport Layer Protocols (cont)

QUIC

- Transport over UDP
- Multi-streaming:
 - Every stream is a **reliable** bidirectional bytestream
 - Multiplexed streams between **two endpoints**
 - Stream prioritization
- Flow-control and congestion control very similar to TCP
- Endpoints use Explicit Congestion Notification (ECN)

SCTP

- Basically, a TCP++
- Multi-streaming
- Unordered delivery is possible
- Flow control and congestion control similar to TCP

Transport Layer Protocols (cont)

What are the limitations of SCTP and QUIC?

- E2E communication: multiple flows (streams) of the same application may connect more than two end-points
- A blind congestion control
- No support from the network: the network knows better about its state
 - ➔ Can better locate and manage congestion
 - ➔ Predict and detect more efficiently congestions/failures/problems...
 - ➔ Can retransmit faster
 - ➔ Can provide better guarantees in terms of delay and packet loss

Network Protocol Stack/Functions

Transport Assistant (cont)

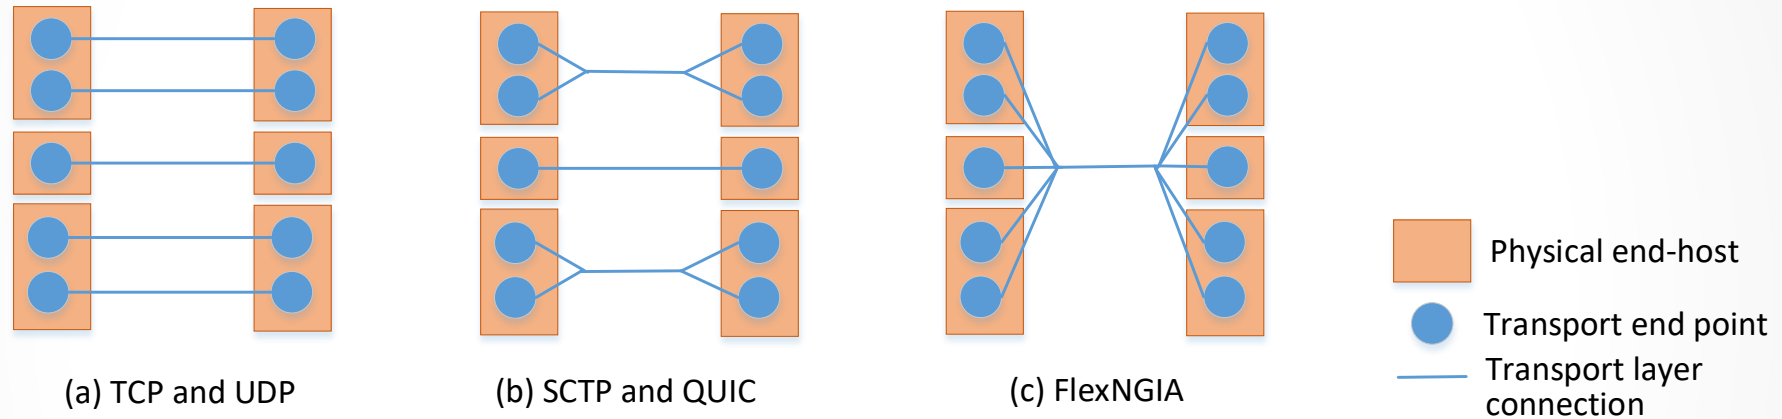


Illustration of how one single network application might be seen at the transport Layer

- Transport Assistants manage all these flows while taking into account that they all belong to the same application
- TAs monitor these flows, divide the total bandwidth allocated for the application among them.